

INTRODUCTION TO LINUX

MARCEL JAR

CONTENTS

I	INTRODUCTION TO COMMAND LINE INTERFACES	1
1	WHAT IS AN OPERATING SYSTEM	2
1.1	Operating System Components	2
1.2	Computer Architecture Basics	4
1.2.1	CPU	4
1.2.2	RAM	4
1.2.3	HD	5
1.2.4	Peripherals	5
	Exercises	5
2	A BRIEF HISTORY OF LINUX	6
2.1	Unix	6
2.2	GNU	7
2.3	Linux	8
3	LINUX DISTRIBUTIONS	10
3.1	Distribution Components	10
3.1.1	Desktop Environment	10
3.1.2	Software Management	12
3.1.3	Default Applications	12
3.2	Popular Distributions	12
3.2.1	Ubuntu, Mint, Kubuntu	13
3.2.2	Debian	13
3.2.3	SUSE, OpenSUSE	13
3.2.4	Red Hat, Fedora, CentOS	13
	Exercises	14
II	INTRODUCTION TO COMMAND LINE INTERFACES	15
4	COMMAND LINE INTERFACE, TERMINAL, AND SHELL	16
4.1	Terminology	17
4.2	Shells	18
	Exercises	18
5	BASIC SHELL COMMANDS I	19
5.1	Accessing a Terminal Emulator	19
5.2	Terminal Basics	19
5.3	date	20
5.4	whoami	20
5.5	pwd	21
5.6	ls	21
5.7	tree	22
5.8	cd	23
5.9	Relative and Absolute Paths	23
5.9.1	Relative Path	24
5.9.2	Absolute Path	25

5.10	clear	25
5.11	Command History	26
	Exercises	26
6	BASIC SHELL COMMANDS II	28
6.1	mkdir	28
6.2	rmdir	29
6.3	touch	29
6.4	rm	30
6.5	mv	31
6.5.1	Moving files and folders accross directories	31
6.5.2	Renaming files and folders	32
6.6	cp	32
6.6.1	Copying files within the working directory	33
6.6.2	Copying files to other directories	33
	Exercises	34
7	GETTING HELP	36
7.1	help	36
7.2	man	37
7.2.1	Sections	38
7.3	whatis	39
7.4	apropos	39
7.5	info	39
	Exercises	40
8	READING AND EDITING TEXT FILES FROM THE SHELL	43
8.1	Reading text files	43
8.1.1	more	43
8.1.2	Less	44
8.2	Editing Text Files	46
8.2.1	nano	46
8.2.2	vi and vim	47
	Exercises	52
III	ADVANCED COMMAND LINE INTERFACES	53
9	LINUX FILE SYSTEMS	54
9.1	File System: Data Storage Format	54
9.1.1	Linux File System Evolution	54
9.1.2	Linux File System Format	55
9.1.3	stat	55
9.1.4	Directories in Linux	56
9.1.5	Accessing data and metadata	57
9.1.6	Performing actions on files, directories, and inodes	57
9.1.7	shred	58
9.2	Directories Hierarchy	59
	Exercises	61
10	FILE LINKS	63

10.1	Hard Links	63
10.1.1	ln	63
10.1.2	Hard Link Properties	66
10.1.3	Hard link limitations	66
10.1.4	Hard Link usage	66
10.2	Soft Links	67
10.2.1	ln -s	67
10.2.2	Soft link properties	70
	Exercises	70
11	FILE GLOBBING: USING WILDCARDS	72
11.1	star - *	72
11.2	Question Mark - ?	74
11.3	Square Brackets - []	75
11.4	Escaping special characters	75
	Exercises	76
	Exercises	76
12	GREP AND REGULAR EXPRESSIONS	77
12.1	grep	77
12.1.1	grep options	78
12.2	Regular Expressions	78
12.2.1	Examples of grep commands using regex	79
	Exercises	81
13	BASIC FILTERS	82
13.1	cat	82
13.2	head	83
13.3	tail	84
13.4	tac	84
13.5	sort	85
13.6	cut	87
13.7	sed	88
	Exercises	89
14	AWK, RENAME, AND FIND	91
14.1	awk	91
14.2	rename	93
14.3	find	94
	Exercises	96
15	REDIRECTION AND PIPING	97
15.1	Redirection	97
15.1.1	Redirection Syntax	98
15.2	Piping	99
15.2.1	Piping Application Scenario	100
15.2.2	Piping Syntax	101
15.2.3	Piping examples	101
15.2.4	Piping and Redirection	102
15.2.5	tee	102

IV	SCRIPTING	103
16	INTRODUCTION TO SCRIPTING	104
16.1	Hello World Script	104
16.1.1	shebang line	104
16.1.2	Comments	105
16.1.3	echo	105
16.2	Using the terminal to run scripts	106
16.2.1	PATH variable	106
16.2.2	Granting permission to execute a Script	107
16.3	Variables and Environments	108
16.3.1	source	108
16.4	Gathering user's input	110
16.4.1	read	110
16.4.2	Passing arguments to a script	110
	Exercises	111

LIST OF FIGURES

Figure 1.1	Basic computer architecture.	4
Figure 4.1	Notepad application for Windows 7.	16
Figure 5.1	Directory tree.	24
Figure 5.2	Directory tree for questions 9, 10, and 11.	27
Figure 6.1	Directory tree for questions 4 and 5.	35
Figure 8.1	Vi operational modes and the keys required to change modes.	49
Figure 9.1	Sequence of steps to access a text file.	58
Figure 10.3	Directory tree for questions 2 and 3.	70
Figure 15.1	Visual representation of data streams.	98

LIST OF TABLES

Table 1.1	Windows OS Components.	3
Table 4.1	List of Linux Shells.	18
Table 5.1	Long list information for the seneca.pdf file.	22
Table 7.1	Manual Page Sections	38
Table 7.2	Shortcuts to navigate info pages	40
Table 8.1	Less navigation keys.	45
Table 8.2	Information displayed in vim	48
Table 8.3	Keys to switch to insert mode.	49
Table 8.4	List of some command mode important com- mands.	50
Table 8.5	List of some extended mode important com- mands.	51
Table 9.1	Contents of the /home/marcel directory file.	56
Table 9.2	Basic directory contents according to the FHS	61
Table 12.1	grep options.	78
Table 12.2	Special characters for regex	79
Table 13.1	sort options.	85
Table 14.1	awk patterns.	92
Table 14.2	find search criteria types.	95
Table 15.1	Bash data streams.	97

LISTINGS

Listing 7.1	Manual page for the mkdir command.	42
Listing 8.1	less user interface.	45
Listing 8.2	Nano's user interface.	46
Listing 8.3	vim 's user interface.	48
Listing 13.1	Poem1 text file	83
Listing 13.2	Poem2 text file	83
Listing 13.3	Countries file.	85
Listing 14.1	Car dealership text file.	91
Listing 15.1	Computing Science students information (cs.info).	100
Listing 15.2	Fine Arts students information (fa.info).	100
Listing 16.1	Script for a Hello World example.	104
Listing 16.2	Script for changing the USER variable.	109
Listing 16.3	Script that uses the read keyword.	110
Listing 16.4	Script that takes arguments from the user's script call.	111

ACRONYMS

CLI	Command Line Interface
GUI	Graphical User Interface
UI	User Interface
OS	Operating System
FHS	File system Hierarchy Standard
PC	Personal Computer
GNU	GNU is Not Unix
FSF	Free Software Foundation
CPU	Central Processing Unit
RAM	Random Access Memory
HD	Hard Disk
SSD	Solid State Device
regex	Regular Expression

CSV Comma-separated values

Part I

INTRODUCTION TO COMMAND LINE INTERFACES

In the following chapters, we explain what an operating system is, as well as its relationship with the computer hardware and software applications. Following, we present a brief overview of Linux's history. Finally, we discuss what Linux distributions are, and we comment on some of the defining characteristics of some of the most widely used distributions.

WHAT IS AN OPERATING SYSTEM

To understand Linux Operating Systems, it is necessary first to understand what an Operating System (OS) is. In a nutshell, an operating system is a fundamental piece software that manages both hardware and software resources, while also providing users with a set of core utilities.

1.1 OPERATING SYSTEM COMPONENTS

An OS is normally divided in many components in order to make the different parts of a computer work together. Amongst those components, it is worth to cite:

KERNEL The kernel is the most crucial part of any OS. It serves as an intermediate between software applications and the hardware, having complete control over everything that happens in the system. It controls which software processes are running at any given time, how much memory each process is assigned to, and also provide a gateway between these processes and hardware such as printers, network cards, keyboards, etc.

USER INTERFACE In order to be useful, an OS needs to provide users with a way to access the data stored in it, as well as run different types of applications. In the past, most OSs had only command line interfaces. However, most modern systems nowadays provide users with increasingly easier to use graphical user interfaces.

FILE SYSTEM The responsibility to store and retrieve data into/from memory devices, as well as ensure that it hasn't been corrupted, belongs to the OS.

SECURITY Any respectable OS must provide methods for users to control access to the system itself, as well as to control access to individual files and processes.

CORE UTILITIES The OS is also expected to provide the user with basic utilities to perform things such as: find files, edit text files, control system settings, control processes, etc.

Probably, the most famous OS is Microsoft Windows, whose version 10 now comes installed by default in a large number of laptops and desktops. It is easy to see that Windows has all the components described above, as seen in Table 1.1.

Kernel	The Windows kernel contains a large set of device drivers that allows it to interact with a very wide array of hardware components, from external keyboards to virtual reality glasses. It is able to run multiple applications at once, and allow users to control which applications get more priority.
User Interface	The graphical interface for Windows is classically composed of a Start Button, a Taskbar, as well as “Windows,” where each “Window” normally correspond to an specific application.
File System	Windows is normally installed using an NTFS file system that ensures the reliability of the stored data.
Security	Windows provides control access to the system, as you normally need to enter a password to access it. Also, some versions of Windows allow users to set specif access rights to individual files.
Core Utilities	Windows comes with multiple core utilities such as: File Explorer to retrieve any file in its memory, Notepad to edit text files, Control Panel to modify systems settings, etc.

Table 1.1: Windows OS Components.

1.2 COMPUTER ARCHITECTURE BASICS

In order to better appreciate the role the kernel, it is important to understand the architecture of modern computers. In Figure 1.1, you can see how the kernel fits within a basic computer architecture. In it, it is clear that the kernel acts as a bridge between applications (at the top), and hardware components (at the bottom). Note that the kernel also provides a bridge between different applications, as well as different hardware components.




Figure 1.1: Basic computer architecture.

In what follows, we provide a brief introduction of some of the most important hardware components of a modern computer.

1.2.1 CPU

The Central Processing Unit, **CPU**, is the computer's "brain". It is responsible for arithmetic and logical operations, as well as control and input/output (I/O) operations. Whenever an application needs to perform an operation, it requests the kernel to send instructions to the **CPU**. After the **CPU** processes the given instructions, it sends outputs back to kernel which, in turn, sends it back to the application.

A single-core computer can only execute one instruction at the time*. When multiple applications are open at the same time, it is the job of the kernel to control access to the **CPU**. The kernel accomplishes it by having applications taking turns in sending instructions to the **CPU**. The kernel also ensures that applications with higher priority levels will get more turns to access the **CPU** than applications with lower priority levels.

**Dual core computers can execute two, and quad-core computers can execute four instructions at time.*

1.2.2 RAM

Random Access Memory, **RAM**, is a type of memory which provides very fast access to retrieve and edit the data stored in it. When an

application starts, the kernel normally loads its contents into **RAM**, together with any data it requires during start-up. Once the application closes, it is the job of the kernel to free **RAM** in order to be used by other applications.

It is also the job of the kernel to ensure that, when multiple applications are open, one application cannot access or overwrite memory addresses used by another application.

1.2.3 HD

The hard disk, **HD**, is where all data is ultimately stored in a computer. Historically, computers have always used magnetic hard-disks for this purpose. However, some new laptops use solid-state drives **SSD***. Both **HDs** and **SSDs** can hold data even without a power source. However, they require a power source to retrieve the data.

All your personal files, as well as application files, and even system files are stored in the **HD**. Whenever an user starts an application, the kernel retrieves the application data in the **HD** and loads it into the **RAM**. This process is the reason why some applications take a few seconds to start*. When an application opens a personal file, the file is retrieved from the **HD** and stored in the **RAM**. When an user creates a new file, or edit an existing file, the kernel is responsible for transferring it from the **RAM** into the **HD**.

**Which are normally faster, but more expensive than magnetic devices.*

This is also why computers equipped with **SSD have faster application boot times.*

1.2.4 Peripherals

Modern computers are normally connected to multiple devices that can request data from the computer, or send data to it. As some examples, we have wireless adaptors, keyboards, computer screens, printers, etc. It is the kernel, using device drivers, that provides a bridge that allow applications to interact with these devices.

EXERCISES

1. With your own words, explain what an operating system is.
2. With your own words, explain what is the role of a kernel.
3. What is the role of the CPU?
4. Why do most computers have two types of memory? I.e., why do computers have **RAM** memory as well as **HD** (or **SSD**) memory?

A BRIEF HISTORY OF LINUX

Most operating systems, such as Microsoft Windows or Apple OS X are well-defined products from a particular company. As a result, all users of one of these systems have exactly the same OS installed on their computers*.

When talking about Linux, quite the opposite is true. There is a myriad of versions of Linux out there, and each one caters for a different type of public or need. There are versions of Linux tailored to run on low-power devices with little amount of memory, such as Raspberry Pis, and versions of Linux running on the world fastest super computers. There are versions of Linux with only a command line interface used to control industrial equipment, and versions of Linux with very friendly graphical user interfaces for phones, tablets, and laptops. There are even versions of Linux powering servers that control most of the traffic of information in the World Wide Web.

In fact, due to its flexibility, and contrary to popular belief, Linux OSs are the most popular operating systems in the world. Even though, in many cases, users are not even aware they are using Linux. For example, Linux was the basis upon which Google's Android OS was developed. They are also the OS of choice for the vast majority of smart devices such as TVs, thermostats, wireless routers, etc. Also, they dominate the server market, with major corporations such as Google and Facebook currently using* Linux servers to run their search engines and host their social media data, respectively.

**Some releases of Microsoft Windows may have different editions such as Server, Home, or Professional.*

**As of 2016.*

To understand how Linux became so popular, and why there are so many different versions of it, it is crucial to learn from where it came from. In what follows, we present a brief overview of its history.

2.1 UNIX

The history of Linux, as well as that of most operating systems, starts in 1969 when Ken Thompson, Dennis Ritchie, and other scientists from AT&T Bell Labs released the first version of the Unix OS.

During the late seventies and early eighties, the popularity of Unix grew tremendously amongst academia and corporations. Part of its popularity had to do with its visionary design choices, such as:

- It was written using the C language*, which made it easier to port it to multiple devices.
- It allowed multiple users to use multiple applications concurrently (at the same time). This was very important back then,

**First it was written in Assembly, but later it was rewritten in C.*

when many universities and corporations had hundreds of users, but only a handful of mainframe computers.

- It had an easy-to-use modular design. Instead of relying on complex tools to perform complex operations, it relied on simple tools that could easily be combined to perform complex operations. See the Unix philosophy box below.

Another reason behind the surge in Unix's popularity was its price tag. Due to antitrust restrictions, AT&T Bell labs could not sell products that weren't specifically targeted for telecommunications. As a matter of fact, they were required to provide the Unix OS source code, free of charge, to anyone who asked for it. As word of mouth had been quite positive, many Universities and corporations required the Unix OS source code and started using it. Also, given that they had access to the source code, they improved it by removing bugs, porting it to more systems, and adding more features to it. In fact, a lot of tools used in Unix to this day were written by students, professors, and developers, and then made freely available.

Unix Philosophy

The Unix philosophy is not a formal design method nor an academic driven methodology. Instead, it is a bottom-up, pragmatic, and grounded in experience way of designing operating systems and software in general.

It was summarized by Peter H. Salus, autor of *A Quarter-Century of Unix*, as a set of three simple rules:

- Write programs that do one thing and do it well.
- Write programs to work together.
- Write programs to handle text streams, because that is a universal interface.

2.2 GNU

In the early eighties, the computer industry grew into a multi-billion dollar business. Thus, as an attempt to increase their profits, some companies started to market their own modified versions of Unix, normally called Unix-like OSs, as closed-source OSs*.

All of a sudden, the Unix world had changed. It went from an open community in which everyone had access to use it, change it as they pleased, and share their results, to one where people and companies had to pay thousand of dollars to use an OS which they could only treat as a black box.

*Perhaps the most famous closed-source Unix-like systems today are Apple's OS X (laptops) and iOS (iPhones and iPads).

Many people, as one would expect, were very displeased by this paradigm shift for several reasons. One of the reasons was that they didn't think it was fair to have to pay to use a tool they had helped to develop. Another reason was that, as different companies were changing their Unix-like OSs in different ways, and the source code was been kept confidential, tools developed in one Unix-like system would often not work properly in other Unix-like system.

One of the more famous voices against this paradigm shift was Robert Stallman. In order to try to re-establish a culture of collaboration, as it was the case during the early days of Unix, he created the Free Software Foundation (FSF) to support a project called GNU*, with the goal of re-creating Unix using an open-source model. His efforts were quite succesful, as many important UNIX tools that were rewritten for the GNU project are still used to this day. Also, the license models championed by the GNU Project, which allowed anyone to use, modify, and redistribute source code, fostered hundreds of other projects. Nevertheless, there was one glaring missing piece in their quest to create an Unix-like OS built entirely using open source code. Their progress in re-creating the Unix kernel was going very slow.

**A recursive acronym meaning GNU is not Unix.*

2.3 LINUX

In 1991, a Finnish student from the University of Helsinki called Linus Torvalds started working in his own re-implementation of the Unix kernel. More specifically, he used an Unix-like system called MINIX as a basis for his own system.

The source code for the resulting kernel, which was dubbed Linux*, was made freely available, together with a call for comments and suggestions for improvement. Because a kernel by itself is not very useful, Torvalds also published a list of GNU tools for Linux users in order to have a functional OS.

**A combination of the words Linus and Unix.*

Many computing enthusiasts started using, and contributing to the further development of an open source OS consisting of a Linux kernel and GNU tools. Over time, the name Linux started to be commonly used to denote the entire OS, as opposed to just the kernel*.

**Some people refer to it as GNU/Linux, though.*

At first, Linux users had to download and compile the kernel and all required tools separately. Hence, Linux acquired a reputation of being difficult to use. However, over time, different developers started to create packages containing the kernel and sets of tools in order to make it easier for other users to install and use it. This resulted in the rise of a number of different Linux distributions that we discuss in our next chapter.

The Rise of MS Windows

From the middle to late eighties, while Unix was gaining traction in academia and high-tech corporations, a competing OS from Microsoft, called DOS, together with its graphical user interface, called Windows, was making great strides in the households and offices market.

The reasons behind the dominance that Microsoft still holds on these markets are plenty, but it is worth to cite a few:

- A deal with IBM in the eighties guaranteed that Microsoft products would come installed by default in all IBM personal computers PCs.
- These PCs found a sweet spot in the computer market. They were significantly cheaper than the more powerful Apple computers at the time, but they were still powerful enough to run the types of applications most users needed at the time. This led to IBM PCs, which used Microsoft Windows, to dominate the computer market by the early nineties.
- They were considered easier to use, as they were focused on a Graphical User Interface (GUI), as opposed to Command Line Interfaces (CLI) common in Unix-like systems at the time.
- Microsoft cemented their OS dominance with Windows-exclusive applications that became ubiquitous in offices and households such as: Word, Excel, and Power Point, which would later be bundled into Microsoft Office.

LINUX DISTRIBUTIONS

As mentioned in the previous chapter, in the beginning of its history, Linux wasn't an OS that catered for unexperienced users. For starters, in order to install it, users needed to download the source code for the Linux kernel, as well the source code for GNU tools, and compile them in their target systems.

In order to remove this technical entrance barrier, some Linux users started to compile the kernel's source code, as well as the source code for a number of GNU tools, for popular computer models. The resulting binaries were then combined into single packages, normally called distributions, which were made available for other users to download. This allowed Linux OSs to be installed without requiring the technical expertise to compile the source code.

Given the open nature of Linux and GNU, anyone could package different sets of tools together in order to create their own distribution. That is exactly what happened. Over time, more and more users packaged their distributions with different sets of components*.

**It is important to note that, over time, many tools that are not part of the GNU project were added to different distributions.*

3.1 DISTRIBUTION COMPONENTS

The one piece of source code that defines an operating system as a Linux OS is the Linux Kernel. Apart from it, the other components can vary tremendously from one distribution to another. In what follows, we describe some of the most important components of Linux distributions.

3.1.1 Desktop Environment

The most noticeable aspect of a Linux distribution is its Desktop Environment. A Desktop Environment provides a graphical interface for users to access different applications, files, as well as control the system settings. Multiple desktop environments were created over the years, with different design proposals, and tackling different issues. A small list of popular desktop environments is provided in what follows. We also provide a few examples of popular distributions that use the discussed environments by default*.

**Many Linux Distributions allow users to switch their Desktop Environment from its default.*

KDE The K Desktop Environment, **KDE**, was the first popular desktop environment created for Linux OSs that is still in use today. It follows a design proposal similar to that of the classical Windows OS. I.e., it provides users with a start button from where

the users can start different applications, a taskbar, and a Desktop where the user can drag and drop files and shortcuts. **KDE** provides many extras bells and whistles that are not available in Windows though, such as multiple workspaces, Desktop Folders, etc. *Ex: OpenSUSE*

GNOME **KDE** required tools that were not open source*. As a result, members of the **GNU** project started to work towards creating a Desktop Environment composed entirely of open source code. The **GNU** Network Object Model Environment, **GNOME**, was the result of this initiative. The **GNOME** Desktop Environment was built with a focus in productivity and provides a set of accessibility tools for users with disabilities. It is currently in its third version (**GNOME 3**). *Ex: Debian, Red Hat, Fedora, CentOS*

**More specifically the QT framework.*

UNITY Unity is a Desktop Environment specifically designed for the Ubuntu distribution. It was designed to make a more efficient use of smaller screens. One of its most noticeable aspects is its launcher bar at the left of the screen. The launcher, as the name suggests can launch different applications, switch between open applications, as well as launch the **dash**, an overlay that allows the user to search quickly for applications, files, folders, etc., both locally and remotely. *Ex: Ubuntu*

OTHERS There are many other desktop environment available that cater for different types of users. Some of these environment were designed for low-power systems, such as **Xfce** and **LXDE**. Other environments are forks of the **GNOME 2*** with added features, for example **Cinnamon** and **MATE**.

**See the GNOME 3 controversy box below.*

GNOME 3 controversy

The third version of **GNOME**, released in 2011, departed significantly from the previous desktop metaphor used in its second version. In this new version, users were no longer provided with a start button and a taskbar that are always visible. Instead, it requires users to constantly switch to an overview mode in which a series of elements such as a dash, a search bar, a list of current workspaces, etc. are available.

Many users complained about such abrupt change in their desktop environment. This lead some users to create other desktop environments based (forked) on **GNOME 2**, such as **Cinnamon** and **MATE**. These environments aim at keeping **GNOME 2**'s desktop metaphor, while at the same time, keeping it up to date with the most modern Linux technology.

3.1.2 Software Management

Virtually all smart phones come with app management tools. These tools are used to download new apps, as well as to keep them up to date. For example, the **App Store** fulfills this role on iOS devices, whereas **Google Play** does the same for Android devices.

Similarly, most Linux distributions come with software management tools which allow users to download and maintain software in their systems. Amongst the most popular package management tools are **apt*** (**Debian**, **Ubuntu**, etc.), and **yum** (**Fedora**, **Red Hat**, etc).

**We cover Software Management on Chapter XXX.*

Most distributions also maintain their own online databases with a list of certified and up-to-date software packages (applications). These databases, called repositories, allow users to quickly retrieve and install new applications. They also allows the system to check which applications are up to date, by comparing the version of the installed application to the latest version in the repositories.

3.1.3 Default Applications

Most Linux Distributions nowadays come with basic terminal **GNU** tools*, **Mozilla Firefox** installed as its browser, and an office application suite called **LibreOffice**, which presents great alternatives for Microsoft Office programs such as Microsoft Word, Microsoft Excel, etc. However, apart from these two applications, different distributions may come with different sets of applications that are installed by default.

**These tools are extensively covered in this book.*

For example, some distributions come with **Mozilla Thunderbird** pre-installed and set as its default email client, while other distributions may come with **Evolution Email**, instead. Some distributions might not come with any pre-installed email client.

Some specialized distributions come with a set of applications for specific purposes. For example, **Kali Linux** is a distribution designed for security experts. Hence, it comes with a set of penetration testing tools. As another example, **Scientific Linux** is a distribution that comes with a set of tools for mathematical modelling, statistical inferencing, and data analysis.

3.2 POPULAR DISTRIBUTIONS

There are, literally, hundreds of Linux distributions available today. This fragmentation comes as a direct result of the open nature of Linux development. Anyone can modify an existing open source distribution to create a new one. This process is normally called forking.

While it is impossible to list all available distributions, it is worth to take a look at some of the most popular ones

3.2.1 *Ubuntu, Mint, Kubuntu*

Ubuntu is, nowadays, the most popular Linux distribution for desktops by a large margin. It has a clear focus on usability and comes with a full gamma of applications to allow users to start being productive without having to make any changes to the system.

Due to its popularity, **Ubuntu** has been forked into a number of new distributions, such as **Linux Mint**, and **Kubuntu**. These distributions mostly differ from **Ubuntu** by presenting different Desktop Environments by default*.

All the examples in the remaining chapters of this book were performed on an Ubuntu OS. However, unless specifically stated otherwise, they should work on any other Linux distribution.

**MATE and Cinamon for Linux Mint, and KDE for Kubuntu.*

3.2.2 *Debian*

Debian is one of the oldest and most succesful Linux Distributions. Its first release dates back to 1993, and so far it has undergone eight major releases*.

The reason behind **Debian** only having had eight major releases is due to its focus on stability. **Debian** development is conducted using three branches: *unstable*, *testing*, and *stable*. Any changes are first tried at the *unstable* branch. Then, after extensive testing by multiple users at the *testing* branch, the change finally finds its way to the *stable* branch.

Due to its focus on stability, **Debian** is one of the favorite distributions for servers. Also, many other Linux distributions are derived from one of the three Debian branches, including **Ubuntu**.

**Each Debian release is named after a characer from the Toy Story movies.*

3.2.3 *SUSE, OpenSUSE*

SUSE is a Linux distribution aimed at enterprises. In contrast to most other Linux Distributions, it cannot be downloaded for free. It needs to be purchased together with an online support plan.

SUSE is based on the **OpenSUSE** distribution, which as the name suggests, can be downloaded and used for free. Novell, the company the owns **SUSE**, refines and enhances **OpenSUSE** to create **SUSE OSs** for specific enterprise goals, such as data center deployments, server deployments, business desktops, etc.

3.2.4 *Red Hat, Fedora, CentOS*

Like **SUSE**, **Red Hat** is an enterprise Linux distribution* that needs to be purchased together with an online support plan. It is mantained by Red Hat, Inc.

**Its full name is Red Hat Enterprise Linux.*

Also, just like **SUSE** is based upon **OpenSUSE**, **Red Hat** is based on **Fedora OS**, a popular Linux distribution sponsored by Red Hat, Inc., but made publicly available for free*. **Fedora OS** is known for its fast development pace, where multiple releases can happen in a year.

**Its source code is also publicly available.*

Due to its leading position in enterprise Linux, many **Red Hat** clones have appeared through the years. The most popular of these clones, **CentOS**, is sponsored by no one less than Red Hat, Inc itself. The main difference between these two distributions is that **CentOS** is an open source project. Hence, it can be obtained and modified free of charge and it comes with no official support.

EXERCISES

1. With your own words, explain what is a Linux distribution.
2. With your own words, explain what is a Desktop Environment.
3. List at least two differences between a Fedora Distribution, and an Ubuntu Distribution.

Part II

INTRODUCTION TO COMMAND LINE INTERFACES

In the following chapters, we introduce the concept of command line interfaces, and explain how they can be accessed, as well as their relationship with the **shell**. Following, we present some basic **shell** commands together with methods to get help about these commands. Finally, we end this part by introducing a series of tools that can be used to read and edit text files using a command line interface.

COMMAND LINE INTERFACE, TERMINAL, AND SHELL

Most User Interfaces (UI) rely on visual cues to guide the user. For example, on Microsoft Windows, the user can use the mouse (or the finger in a touchscreen device) to click on an icon representing a folder to see its contents. Another example would be that, by clicking on a x symbol at the top right edge* of an application, the application is closed. As yet another example, drop-down menus with items such as File, Edit, and About, are commonplace, as seen in Figure 4.1.

**In Apple and Linux systems, this symbol might be on the top left edge.*




Figure 4.1: Notepad application for Windows 7.

A UI that relies of visual cues, as the ones described above, is called a Graphical User Interface (GUI). The invention of such interfaces played a major role in the popularization of computers in the 80s. This is due to the fact that most users are more comfortable using such interfaces, as opposed to Command Line Interfaces (CLI) that have been around since the late 60s.

In a CLI, as opposed to an intuitive graphical interface, the user is presented with a **shell*** prompt in which it can write commands to be interpreted by the **shell**. You can see an example below of a CLI in which a user entered the command **ls** to list all files and folders in the working directory*.

```
marcel@dell:~$ ls
Documents      Downloads      Pictures
Music          Video         seneca.pdf
marcel@dell:~$
```

**Shells are explained in Sections 4.1 and 4.2.*

In this book, the words **folder and **directory** are used interchangeably.*

In order to properly use a **CLI**, the user needs to learn a series of commands that the shell understands, as well as the syntax of these commands. Hence, it should come to no surprise that most regular users prefer to use **GUIs** which do not incur in such steep learning curves.

If you are reading this book, however, I expect you not to be a regular user. In fact, I expect you to be starting a career as systems analyst, a software developer, or other positions in the IT industry. For people like you, having a clear understanding of how to use a **CLI** presents a number of advantages:

- There are tasks that can only be achieved by using a **CLI** because no **GUI** has been created to perform them.
- Some tasks can be performed, using a **CLI**, in a fraction of the time it would take using a **GUI** tool.
- **CLIs** can be accessed remotely in a fast and straightforward fashion, allowing you to control multiple systems from a single computer
- Users can quickly create scripts that run on **CLIs** to automate repetitive tasks such as: adding new users to the system, backing up data, updating the system, etc.
- The steps required to perform some actions using **GUIs** can vary a lot depending on which **OS**, or Linux distribution, you are using. On the other hand, they are normally identical using a **CLI**. In fact, the same commands* can be used in all Linux distributions, Apple computers, and since 2016, even in Windows systems.

**With a few notable exceptions.*

Most Linux distributions have a **GUI** that allow regular users to achieve daily tasks such as web browsing, email, or writing documents. However, all Linux distributions also come with a terminal emulator that allow super users to do much more using a **CLI**.

4.1 TERMINOLOGY

Before moving forward to learn how to use a Linux **CLI**, It is important to notice the difference between three concepts that are often used interchangeably:

CLI The **CLI** is just the technical name of the method by which we interact with a computer by means of written instructions.

TERMINAL All major Linux distributions have a **GUI**. However, they normally also provide you with a terminal, or more precisely a terminal emulator. A terminal emulator is nothing else than a program that provides us with a **CLI** to interact with the **shell**.

SHELL The **shell** is a software that reads keyboard commands and passes them to the **OS** to carry out. Such commands could be used to perform simple tasks like listing the files in a particular folder, adding users to the system, creating folders, or complex tasks such as upgrading the **OS**, installing applications, etc.

4.2 SHELLS

As mentioned before, a **shell** is a software that interprets commands passed by the user, and passes them to the **OS**. However, just like there are multiple languages such as English, Spanish, or Mandarin, there are multiple shells available. Table 4.1 below presents a description of a few such shells*.

**Many other shells exist, such as dash, ksh, etc.*

sh	The Bourne shell, which has been distributed as the default shell for Unix since 1979.
bash	Bourne Again Shell, an upgraded version of sh , written by the GNU Project, which has become the <i>de facto</i> standard Linux shell.
csh or tcsh	A Unix shell using a syntax similar to the C programming language. tcsh is identical to csh , with the addition of some extra features such as command-line completion.

Table 4.1: List of Linux Shells.

In this book, we will cover exclusively the **bash shell**. This is due to the fact that **bash** is the default shell for the most widely used Linux distributions. In fact, it is hard to find a Linux distribution that does not come with a **bash shell**. That said, after learning how to use one shell, it is much easier to learn other shells, as they share many concepts and syntax.

EXERCISES

1. Explain, with your own words, what is a Command Line Interface (**CLI**).
2. Explain, with your own words, what is a Graphical User Interface (**GUI**).
3. What is the relationship between a terminal emulator, and a shell?
4. Which commands are used to list all files in the working directory using the following shells: sh, bash, csh, dash, and ksh?

BASIC SHELL COMMANDS I

In this chapter we will cover how to get access to a terminal emulator, as well as some very basic **bash** commands.

5.1 ACCESSING A TERMINAL EMULATOR

Once a user logs into a Linux OS with a GUI, there are multiple ways to get access to a **terminal emulator**. Some methods work only on a specific distribution, whereas some methods work on most distributions. In what follows, a list of methods to access terminal emulators are presented for some of the most popular desktop environments*. The most common Linux distribution for each desktop environment is presented in the margin notes.

**Desktop environments are covered in Appendix XXX.*

GNOME3 Press the super (windows) button to call the list of applications, and then type **terminal** in the search field, or search for the terminal in the provided applications list.

Debian, Red Hat, Fedora, CentOS, Kali

KDE Click on the start button at the lower left, and then click on accessories, and finally on **terminal**.

OpenSUSE, Kubuntu, Slackware

UNITY Press the super (windows) button to call the dash (or conversely click on the dash button at the upper left), and then type **terminal**.

Ubuntu

In all of these desktop environments, the terminal emulator can also be started by pressing **ctrl+alt+T**.

5.2 TERMINAL BASICS

Once the terminal emulator starts, the user is presented with an application displaying a blinking cursor at a **shell prompt**, just like shown below.

```
marcel@dell:~$
```

This prompt contains important information about the current parameters of the shell, namely: the username (**marcel**), system name (**dell**), and working directory (represented by a tilde [~]).*. A brief description of these parameters follows:

**The \$ separates command prompts from these current parameters.*

USERNAME Linux systems, as any Unix-like system, was designed to allow multiple users to access the system. When a terminal emulator starts, it assumes the user to be the same one that logged

in into the GUI system. in Chapter XX, we will show how the **su** command can be used to switch users. In the example above, the username is **marcel**.

SYSTEM NAME During the installation of Linux systems, users are required to give your system (local machine) a name. When a terminal emulator starts, it assumes that you are using the local machine. In chapter XXX we will see how to log into remote machines using the **ssh** command. In the example above, my local machine is called is **dell**.

WORKING DIRECTORY Imagine that you issue a command to create a file. In which directory will this file be created? The answer is: in the current **working directory**. When a terminal emulator starts, it assumes by default that you are at the home folder (directory) of the user currently logged in. Hence, whichever actions you perform will, by default, affect this directory. In Section 5.8, we will show how to use the **cd** command to change directories. In the example above, this home folder, which is located at **/home/marcel** is represented by a tilde (~).

In what follows, a number of basic **bash*** commands are presented.

5.3 **date**

To ask what day is today, you can type **date** in the terminal emulator and press enter. The shell will verify the date with the **OS**, display it in the terminal, and start a new prompt line ready to receive more commands, as shown below*. Note that the terminal also shows the current time and the time zone (Eastern Daylight Time - EDT).

```
marcel@dell:~$ date
Wed May 18 19:20:55 EDT 2016
marcel@dell:~$
```

Note that shell commands are **case-sensitive**. I.e., the shell differentiates between upper-case and lower-case letters. For example, while **date** is a valid command, **Date** and **DATE** are not.

5.4 **whoami**

Not all terminal emulators display the username at each command prompt. Therefore, the shell needs to provide a method to verify the username of the current user. This is accomplished by the **whoami** (who am I) command, as shown below.

```
marcel@dell:~$ whoami
marcel
```

You can easily switch shells by typing the name of the shell you want to switch to. For example, you can type **dash to start using a **dash shell**.*

**For the following commands, we will omit this new prompt line.*

5.5 pwd

As mentioned before, when a terminal emulator is open, it needs a directory to act as a starting point. In most Linux system, this starting point is the user's home folder, located at `/home/username`, represented by a tilde (`~`).

As we will show later, users can move to different directories using the shell. The shell's current location is called **working directory**. This is due to the fact that this is the directory that the shell will, by default, act (work) upon. In order to see in which directory you are currently at, you can type **pwd*** in the terminal emulator and press enter. The shell will then print the working directory in the terminal, as shown below.

**short for print
working directory.*

```
marcel@dell:~$ pwd
/home/marcel
```

5.6 ls

To obtain a list of all files and folders present in the working directory, as demonstrated in Chapter 4, you can use the **ls*** command. The shell will then print a list of all files and folders, as shown below*.

**short for list.*

**Some terminal
emulators use
colours to
distinguish between
files, folders, scripts,
etc.*

```
marcel@dell:~$ ls
Documents      Downloads      Pictures
Music          Video         seneca.pdf
```

The list command, by default, does not show hidden files*. In order to do so, you must add the option **-a** or **-all** to the **ls** command, as shown below:

**Hidden files are
files whose names
starts with a . (dot).*

```
marcel@dell:~$ ls
Documents      Downloads      Pictures
Music          Video         seneca.pdf
marcel@dell:~$ ls -a
.              ..            Documents
Downloads     Pictures     Music
Video        seneca.pdf  .System.conf
```

Note that in the command above, the file **.System.conf** only appears when the **ls** command is issued with the **-a** option. Also, note that two extra directories appear: a folder (**.**) and a folder (**..**). These are not real directories. They are simply hard links to the self (**.**) and parent (**..**) directories.

The list command can also be used to gather information about the files and folders in the current working directory. To do so, the **long list** option **-l** must be invoked, as shown below:

```
marcel@dell:~$ ls -a -l
drwxrwxr-x  3 marcel marcel  4096 Jun 21 23:06 .
```

```

drwxr-xr-x 54 marcel marcel 4096 Jun 21 18:59 ..
drwxrwxr-x 2 marcel marcel 4096 Jun 21 23:06 Documents
drwxrwxr-x 2 marcel marcel 4096 Jun 21 23:06 Downloads
drwxrwxr-x 2 marcel marcel 4096 Jun 21 23:06 Pictures
drwxrwxr-x 2 marcel marcel 4096 Jun 21 23:06 Music
drwxrwxr-x 2 marcel marcel 4096 Jun 21 23:06 Video
-rw-rw-r-- 1 marcel marcel 12238 Jun 29 22:54 seneca.pdf
-rw-rw-r-- 1 marcel marcel 126 Jun 28 20:52 .System.conf

```

This long list provides a number of columns containing information about each file. In Table 5.1, we explain what the information in each column represents using the file **seneca.pdf** as an example.

column	example	meaning
1	-rw-rw-r-	File type and permission sets. These topics are covered in Chapter XX.
2	1	Number of hard links to this file. Links are covered in Chapter XX.
3	marcel	Name of the user owner of the file. The concepts of users and ownership are covered in Chapter XX.
4	marcel	Name of the group owner of the file. The concepts of groups and ownership are covered in Chapter XX.
5	12238	Size of the file in bytes.
6	Jun 29 22:54	Timestamp indicating when the file was edited for the last time.
7	seneca.pdf	File name.

Table 5.1: Long list information for the **seneca.pdf** file.

5.7 **tree**

The **tree** command is somewhat similar to the **ls** command. They both display a list of files and folders in the working directory. However, contrary to **ls**, in which any subfolder is simply listed alongside files, the **tree** command goes inside each subfolder and displays the files contained in them, creating a tree-like structure. See the example below:

```

marcel@dell:Seneca$ tree
.
|-- academic_honesty.pdf
|-- calendar.pdf
|-- OPS105

```

```
| |-- students_list
| |-- grades.xls
|-- SRT311
| |-- students_list
| |-- grades.xls
```

```
2 directories, 6 files
```

```
marcel@dell:Seneca$
```

The **tree** command is not available by default in some Linux distributions. To install it, you must write the command **sudo apt-get install tree**, and enter your user password.

5.8 **cd**

To switch working directories, you need only to use the **cd*** command, followed by the name of the directory you want to go. For example, assuming that the folders shown in the previous command's output do exist, typing **cd Documents** results in:

**short for change directory.*

```
marcel@dell:~$ pwd
/home/marcel
marcel@dell:~$ cd Documents
marcel@dell:Documents$ pwd
/home/marcel/Documents
```

The command **cd** by itself always sends the user back to the user's home folder. Also, the command **cd ..** sends the user to the parent folder of the working directory, and the command **cd -** sends the user back to the previous working directory. The parent folder is the folder hierarchically above the current folder. For example, the folder **/home/marcel** is the parent folder of **/home/marcel/Documents**.

Finally, note that if you enter the name of a directory that does not exist, the shell will return an error message and then it will start a new prompt line ready to receive more commands, as shown next.

```
marcel@dell:~$ cd DOCUMENTS
bash: cd: DOCUMENTS: No such file or directory
marcel@dell:~$
```

5.9 RELATIVE AND ABSOLUTE PATHS

In our previous example using the **cd** command, we switched working directories from one folder, **/home/marcel**, to one of its immediate subfolders, **/home/marcel/Documents**. However, you might face more complex situations in which you have a directory tree as shown in Figure 5.1, and you want to switch the current working directory from any folder directly to any other folder.




Figure 5.1: Directory tree.

For example, given the directory tree shown in Figure 5.1, how can we directly switch working directories for the following cases:

- from **/home** directly to **/home/marcel/Personal**
- from **/home/marcel/Seneca** directly to **/media**
- from **/home/marcel/Seneca** to **/media/Seneca**

The solution to all these cases is: “by providing either a **relative path** or an **absolute path** to the desired folders.” In what follows, we define these two concepts*.

**Note that the concepts of relative and absolute paths also apply to files.*

5.9.1 Relative Path

When a relative path is provided, the shell assumes that the starting point of the path is the current working directory. For example, given that you are currently in **/home**, you can switch your working directory to **/home/marcel/Personal** by issuing the command: **cd marcel/Personal**, as shown in what follows.

```

marcel@dell:home$ pwd
/home
marcel@dell:home$ cd marcel/Personal
marcel@dell:Personal$ pwd
/home/marcel/Personal
  
```


5.9.2 Absolute Path

Absolute paths always start with `/` and provide the complete path for the desired folder or file. For example, given that your working directory is `/home/marcel/Seneca`, you can switch to `/media` by issuing the command: `cd /media`, or you can switch to `/media/Seneca` by issuing the command: `cd /media/Seneca`, as shown in what follows.

**/ denotes the root, of all directories in a Linux system.*

```
marcel@dell:Seneca$ pwd
/home/marcel/Seneca
marcel@dell:Seneca$ cd /media
marcel@dell:media$ pwd
/media

marcel@dell:Seneca$ pwd
/home/marcel/Seneca
marcel@dell:Seneca$ cd /media/Seneca
marcel@dell:Seneca$ pwd
/media/Seneca
```

Relative and absolute paths can be compared to addresses in navigation systems. For example, if you live in Toronto, and you want to go to a place within the city, you can type on your navigation device an address such as 200 King Street. The device will assume that this address is in Toronto-ON, Canada. This would be compared to a relative address, as the starting point is assumed to be the city of Toronto. However, if you are in Toronto and you want to go to 200 King Street in Buffalo-NY, you would need to enter the complete address, including the city, state, and sometimes even the country. This would be compared to an absolute value, as the address is completely defined.

5.10 **clear**

All commands you enter, as well as their outputs, stay in the screen and you are given a new shell prompt at the bottom to keep working with the shell. This can lead to a very polluted screen containing a lot of information that is no longer necessary.

In order to clear the screen from all previous commands and outputs, you simply have to issue the **clear** command. This command will erase from your terminal all information in it and provide you with a command prompt at the top to keep working with the shell.

It is important to note that this command simply clears the terminal display. Any variables* that have been defined, or any modifications to the state of the shell will not be changed by this command.

**Variables are covered in Chapter XXX.*

5.11 COMMAND HISTORY

You can use the up (↑) and down (↓) arrow keys in order to navigate the commands you have previously entered in your shell. For example, by hitting the up arrow key once, you will get the previously entered command. Hitting the up arrow again, will get you the command you entered prior to that one.

EXERCISES

1. Describe two methods to open a terminal emulator on a Ubuntu Linux OS.
2. Given the command prompt below, who is the user currently logged into the shell? Also, what is the name of the local machine? Finally, what is the current working directory?

```
john@lenovo:~$
```

3. What is the relationship between a **terminal emulator**, and a **shell**?
4. Which commands are used to list all files in the working directory using the following shells: bash, csh, dash, and ksh?
5. Which command can be used to show in the terminal what time is it?
6. What is the output of the **ls -l** command? What does the 6th column represents?
7. How can you move back to the latest previously accessed working directory using only one shell command?
8. What happens when you apply the parent folder parameter to the **cd** command twice. In other words, what happens when you issue a **cd ../..** command?
Given the directory tree shown in Figure 5.2 (on next page), answer the following questions:
9. Write a command to switch your working directory from **marcel** to **Seneca** using only one shell command?
10. Write a command to switch your working directory from **media** to **marcel** using only one shell command.
11. Given that your working directory is **media**, can you use a relative path to move to **Seneca**? If yes, can you also use an absolute path? Which one of these options seem more appropriate in this case?

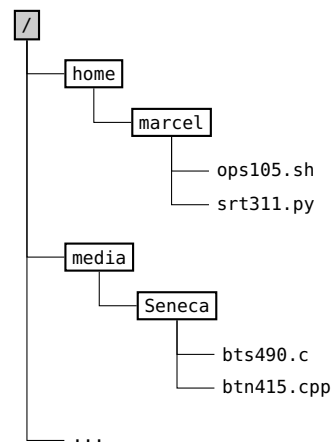


Figure 5.2: Directory tree for questions 9, 10, and 11.

BASIC SHELL COMMANDS II

In our previous chapter, we covered a number of bash commands which could be used to: gather information from the system and user (**date**, **whoami**, **pwd**), get information about files and folders (**ls**), or to change the current working directory (**cd**).

None of the commands could be used to alter files, folders, or change the configurations of the system, though. I.e., they have no lasting effect in the system.

In this chapter, we will cover a series of commands that have lasting effects in files, folders, and possibly the system. The commands we are covering in this chapter can be used to create and remove files and folders, as well as to rename and make copies of them.

6.1 **mkdir**

To create new folders, the **mkdir*** command can be used, followed by the name(s) of the folder(s) to be created. In the example below, a folder called **Samples** is created in the current working directory.

**short for make directory.*

```
marcel@dell:~$ ls
Documents      Downloads      Pictures
Music          Video          seneca.pdf
marcel@dell:~$ mkdir Samples
marcel@dell:~$ ls
Documents      Downloads      Pictures      Samples
Music          Video          seneca.pdf
```

Note that, by default, the new directory is created as a subfolder of the working directory. To create folders in other locations, without leaving the current working directory, you can use the **mkdir** command combined with relative or absolute paths, as shown below:

```
marcel@dell:~$ ls Documents
Seneca      Personal
marcel@dell:~$ mkdir ~/Documents/Samples
marcel@dell:~$ ls Documents
Seneca      Personal      Samples
```

Files and Folders with spaces in their names

In all examples that were given so far in this book, all files and folders do not contain spaces in their names. Issuing a command such as **mkdir My Documents** would create two folders,

one called **My** and another called **Documents**. This is because the shell cannot distinguish between an argument with a space, or two arguments without spaces.

Linux can handle files and folder with spaces, though. In order to do so, whenever you are referring to a file or folder with spaces in its name, each space needs to be preceded by a backslash (\). For example, you can create a folder called **My Documents** by issuing a command `mkdir My\Documents`.

6.2 `rmdir`

The `rmdir`* command, as the name suggests, deletes existing folders. For example, `rmdir Samples` will remove a folder **Samples** from the current working directory. It is important to note that it can only be used to remove empty folders. In case the folder is not empty, the shell returns an error message.

**short for remove directory.*

In the example below, the `rmdir` command is applied to both an empty folder **Samples**, and to a non-empty folder **Pictures**.

```
marcel@dell:~$ ls
Documents      Downloads      Pictures      Samples
Music          Video          seneca.pdf
marcel@dell:~$ rmdir Samples
marcel@dell:~$ ls
Documents      Downloads      Pictures      Music
Video          seneca.pdf
marcel@dell:~$ rmdir Pictures
rmdir: failed to remove 'Pictures': Directory not empty
marcel@dell:~$
```

6.3 `touch`

As shown in Section 5.6, files and folders in Linux have a timestamp indicating when they were edited for the last time. The `touch` command, when applied to an existing file, updates the timestamp of when it was last edited*. It is equivalent to opening the file, saving it, and then closing it. As an example, see the command prompt below:

**Updated timestamps are useful for some backup programs, among other uses.*

```
marcel@dell:~$ ls -l
total 0
-rw-rw-r-- 1 marcel marcel 0 Jul 12 11:31 final_exam.doc
-rw-rw-r-- 1 marcel marcel 0 Jun 21 20:30 introduction.ppt
-rw-rw-r-- 1 marcel marcel 0 May 10 21:41 seneca.pdf
marcel@dell:~$ touch final_exam.doc
marcel@dell:~$ ls -l
```

```
total 0
-rw-rw-r-- 1 marcel marcel 0 Jul 14 22:12 final_exam.doc
-rw-rw-r-- 1 marcel marcel 0 Jun 21 20:30 introduction.ppt
-rw-rw-r-- 1 marcel marcel 0 May 10 21:41 seneca.pdf
```

When applied to non-existing files, the **touch** command creates empty files, as shown below:

```
marcel@dell:~$ ls -l
total 0
-rw-rw-r-- 1 marcel marcel 0 Jul 14 22:12 final_exam.doc
-rw-rw-r-- 1 marcel marcel 0 Jun 21 20:30 introduction.ppt
-rw-rw-r-- 1 marcel marcel 0 May 10 21:41 seneca.pdf
marcel@dell:~$ touch mid_term.doc
marcel@dell:~$ ls -l
total 0
-rw-rw-r-- 1 marcel marcel 0 Jul 14 22:12 final_exam.doc
-rw-rw-r-- 1 marcel marcel 0 Jun 21 20:30 introduction.ppt
-rw-rw-r-- 1 marcel marcel 0 May 10 21:41 seneca.pdf
-rw-rw-r-- 1 marcel marcel 0 Jul 14 22:15 mid_term.doc
```

This second usage of the **touch** command is frequently used for testing purposes*.

tip: you to use **touch and **mkdir** in order to recreate the examples provided in this chapter..*

**short for remove.*

6.4 rm

To delete (remove) files, the **rm*** command can be used, followed by the name(s) of file(s) to be deleted from the working directory. See the example below:

```
marcel@dell:~$ ls
Documents      Downloads      Pictures
Music          Video          seneca.pdf
marcel@dell:~$ rm seneca.pdf
marcel@dell:~$ ls
Documents      Downloads      Pictures
Music          Video
```

The **rm** command can also be used to remove non-empty folders. To do so, we have to apply the **-r** (recursive) option followed by the name of the folder, as shown below.

```
marcel@dell:~$ ls
Documents      Downloads      Pictures
Music          Video          seneca.pdf
marcel@dell:~$ rm -r Pictures
marcel@dell:~$ ls
Documents      Downloads      Music
Video          seneca.pdf
```

Bash Design Goals

It is important to note that the commands described in this chapter don't normally ask for confirmation. For example, when deleting a file using the **rm** command, the **shell** will not ask the user if he/she is sure he/she wants to delete the file. Nor the file will be sent to a recycle bin from which it can be recovered. The file will simply be deleted.

The **shell** performs tasks that might incur in unforeseen consequences, such as accidentally deleting a number of files, in a much more direct way because of its design goals. As opposed to **GUI** interfaces where the goal normally is to provide a platform for inexperienced users to perform basic tasks, a **CLI** is normally designed with the goal of allowing experienced users to quickly perform complex tasks.

6.5 mv

The **mv*** command can be used in two distinct ways. As the name suggests, it can move a file, a folder, or a number of files and folders from one directory to another. However, it can also be used to rename single files or folders. Both uses of this command are described in what follows.

*short for move.

6.5.1 Moving files and folders accross directories

To move a single file from one directory to another, the **mv** command needs two arguments. The first being the name of the file in the working directory to be moved, and the second being the absolute or relative path of the directory we are moving the file to. See the example below:

```
marcel@dell:~$ ls
final_exam.doc  introduction.ppt  seneca.pdf
Folder
marcel@dell:~$ mv final_exam.doc Folder
marcel@dell:~$ ls
introduction.ppt  seneca.pdf  Folder
marcel@dell:~$ ls Folder
final_exam.doc
marcel@dell:~$
```

Note that, in this example, the file **final_exam.doc** disappears from the working directory and appears in the **Folder** directory.

Multiple files can be moved with a single **mv** command. To do so, you first need to provide the names of all files you intend to move as

arguments. Then, you need to provide the relative or absolute path of the directory you want to move them to as the very last argument. This technique is shown in what follows:

```
marcel@dell:~$ ls
final_exam.doc  introduction.ppt  seneca.pdf
Folder
marcel@dell:~$ mv final_exam.doc introduction.ppt Folder
marcel@dell:~$ ls
seneca.pdf  Folder
marcel@dell:~$ ls Folder
final_exam.doc  introduction.ppt
marcel@dell:~$
```

Note that, within this context, this command can also be applied to folders (directories). I.e., you can move an entire folder in your working directory into other folders using the exactly same syntax applied for files in this section.

6.5.2 Renaming files and folders

Imagine the **mv** command being used with two arguments, the first being a file name in the working directory, and the second being a name that doesn't match any directories. In this scenario, the **mv** command will rename the file indicated by the first argument to the name specified in the second*. See the example below:

```
marcel@dell:~$ ls
final_exam.doc  introduction.ppt  seneca.pdf
Folder
marcel@dell:~$ mv final_exam.doc final_exam_fall.doc
marcel@dell:~$ ls
final_exam_fall.doc  introduction.ppt  seneca.pdf
Folder
marcel@dell:~$
```

A command called **rename exists in **bash**. However, it is used to rename multiple files at once using regular expressions. Regular expressions are covered in Chapter XXX.*

In this example, you can see that the file **final_exam.doc** was renamed to **final_exam_fall.doc**.

You can also rename a folder using the same syntax applied for files in this section.

6.6 cp

The **cp*** command can be used to copy files and folders. It can be used in two distinct ways:

**short for copy.*

6.6.1 *Copying files within the working directory*

To create a copy of a file in the same working directory, the **cp** command should be used with two arguments. The first argument should be the name of the file to be copied, while the second argument should be the name of the copy, as shown below:

```
marcel@dell:~$ ls
final_exam.doc  introduction.ppt  seneca.pdf
Folder
marcel@dell:~$ cp final_exam.doc final_exam_fall.doc
marcel@dell:~$ ls
final_exam.doc  final_exam_fall.doc  introduction.ppt
seneca.pdf  Folder
marcel@dell:~$
```

To copy an entire folder in your working directory into another folder also in the working directory, you must use the recursive option **-r**, as shown in the example below:

```
marcel@dell:~$ ls
final_exam.doc  introduction.ppt  seneca.pdf
Folder
marcel@dell:~$ cp -r Folder Folder_Copy
marcel@dell:~$ ls
final_exam.doc  final_exam_fall.doc  introduction.ppt
seneca.pdf  Folder  Folder_Copy
marcel@dell:~$
```

6.6.2 *Copying files to other directories*

To create a copy of a file in the working directory to another directory, the second argument of the **cp** command must be the absolute or relative path of the directory in which you want to place a copy of the file. In this case, the new file will have the same name as the original file. See the example below:

```
marcel@dell:~$ ls
final_exam.doc  introduction.ppt  seneca.pdf
Folder
marcel@dell:~$ cp final_exam.doc Folder
marcel@dell:~$ ls
final_exam.doc  introduction.ppt  seneca.pdf
Folder
marcel@dell:~$ ls Folder
marcel@dell:~$ ls
final_exam.doc
marcel@dell:~$
```

Just as with the previous scenario, to copy an entire folder in the working directory to another directory, you also must use the recursive option **-r**. See the example below:

```
marcel@dell:~$ ls
final_exam.doc  introduction.ppt  seneca.pdf
Folder Music
marcel@dell:~$ cp Music Folder
cp: omitting directory 'Music'
marcel@dell:~$ ls Folder
final_exam.doc  introduction.ppt
marcel@dell:~$ cp -r Music Folder
marcel@dell:~$ ls Folder
final_exam.doc  introduction.ppt  Music
marcel@dell:~$
```

To copy all contents of a folder to another folder, but not the folder itself, you need to add **./** to the end of the name of the folder you are copying, as shown below*:

```
marcel@dell:~$ ls
final_exam.doc  introduction.ppt  seneca.pdf
Folder Music
marcel@dell:~$ cp -r Music/. Folder
marcel@dell:~$ ls Folder
arcade_fire-ready_to_start.mp3  foo_fighters-walk.mp3
marcel@dell:~$
```

You can also use this technique with the **mv command.*

EXERCISES

1. What is the main difference between using the copy **cp**, and the move **mv** commands, when applied to files?
2. How can you delete a non-empty folder called **Archive** that exists in your working directory?
3. Which command can be used to move a file called **README.txt** from your working directory to its parent folder?
4. Which command can be used to create a folder called **Examples** inside a subfolder called **Documentation** that exists in your working directory. See the diagram in Figure 6.1.
5. Still with regards to the diagram in Figure 6.1, how can you move files **Example1.txt** and **Example2.txt** from the **Documentation** folder to its **Examples** subfolder, without changing your working directory.
6. Is it possible to delete multiple empty folders with only one **rmdir** command? If so, how can you delete two empty folders named **Tests** and **Assignments** from your working directory?




Figure 6.1: Directory tree for questions 4 and 5.

7. What happens when you apply the **touch** command to an existing file? Also, what happens when you apply this command to a non-existing file.
8. How can you make copies of the contents of a file called **exams.txt** from your working directory, to a file called **exams.txt** inside a subfolder **Classes** that exists in your working directory?
9. What happens when the first parameter of a **mv** command is a file, and the second parameter is a folder?
10. Which command can be used to copy the contents of a subfolder named **Folder1**, that exists in the current working directory, to another subfolder named **Folder1Copy** in the same working directory. Note that the subfolder **Folder1Copy** doesn't exist prior to the command you need to enter.

GETTING HELP

We have already covered in previous chapters a number of commands which can take multiple options and different numbers of arguments. For example, it was shown in Chapter 5 that the command **ls** can work in three different ways:

- By itself, this command prints the names of all files in the working directory.
- When provided with the name of another folder as an argument, this command prints the names of all files in the specified folder.
- If called with the option **-l***, this command gives all the information about files described on Table 5.1, as opposed to just listing their names.

There are many more options for the **ls command..*

The bash shell has hundreds of commands like **ls** that can take multiple options and different numbers of arguments. Hence, in order to be accessible to non-experts, it needs to provide its users with a way of knowing how to use those commands.

In this chapter we will cover different ways in which users can get help on how to use different commands.

7.1 **help**

By itself, the **help** command will list all **built-in*** commands. If one of the built-in commands is provided as an argument, this command provides a quick description of the provided command, and possibly a list of options. You can see below the output of entering **help pwd** on bash.

**Built-in commands are explained in the Built-in Commands box that follows.*

```
marcel@dell:~$ help pwd
pwd: pwd [-LP]
    Print the name of the current working directory.

Options:
    -L  print the value of $PWD if it names the current
        working directory
    -P  print the physical directory, without any symbolic
        links

By default, 'pwd' behaves as if '-L' were specified.
```

Exit Status:

Returns 0 unless an invalid option is given or the current directory cannot be read.

Note that **help** only covers built-in commands. It does not cover commands implemented as binary files in **/bin** or **/usr/bin**, such as **ls**, **rm**, **mv**, and many others.

Built-in Commands

Most commands in bash, such as **ls**, **rm**, and **touch** are implemented by binary files located in the **/bin** and **/usr/bin** folders. These commands are interpreted in the same way as any other application the shell can run. I.e., the shell asks the kernel to execute it, and after it has finished executing, the shell receives its output.

Built-in commands, on the other hand, are an integral part of the shell. They are not implemented in a separated file. The main reason why some commands are implemented directly inside the shell is because they need to change the state of the shell. For example, the **cd** command is a built-in command because it changes the current working directory. Commands implemented as binaries, such as **ls**, **rm**, and **touch**, cannot change the state of the shell.

Another reason for implementing some commands directly into the shell is because it normally enhances their performance.

7.2 man

Since the inception of Unix, it has become standard practice for the authors of scripts that implement shell commands to provide manual pages for them. This practice was continued by the GNU project when rewriting Unix as an open source project*.

All manual pages follow the same structure show in Listing 7.1 (page 42). I.e., they have the following sections:

Name States the name and purpose of the command

Synopsis Briefly describes the command syntax

Description Describes the command as well as its options

Authors Lists the authors of the script that implements the command

Reporting Bugs Provides a link to a page where bugs can be reported

Copyright States that the code is provided as free software

In fact, the author of the manual pages for many basic commands such as **ls and **rm** is no one less than Robert Stallman, the programmer that started the GNU project, as discussed in Chapter XXX.*

See Also Provides a list of related commands

To access manual pages, you need only to use the **man** command, followed by the name of the command you are trying to get information about.*. See the example below:

```
marcel@dell:~$ man mkdir
```

**Manual pages cover not only commands, but also daemons and config files.*

7.2.1 Sections

Note that the **mkdir** command appears with a **(1)** next to it at the top of Listing 7.1 (see page 42). This number denotes the section of the manual from where the information was retrieved. Some commands have the same name of **deamons** or **config files**. Hence, dividing manual pages in sections makes it possible for users to access the manual page for the right tool. Manual pages are divided in 9 sections, as shown in Table 7.1.

1	Executable programs or shell commands
2	System calls (functions provided by the kernel)
3	Library calls (functions within program libraries)
4	Special files (usually found in /dev)
5	File formats and conventions e.g. /etc/passwd
6	Games
7	Miscellaneous (including macro packages and conventions), e.g. man(7) , groff(7)
8	System administration commands (usually only for root)
9	Kernel routines [Non standard]

Table 7.1: Manual Page Sections

By default, **man COMMAND** retrieves the first occurrence of **COMMAND** in the manual pages. In order to access later occurrences, you need to provide the section as the first argument followed by the name of the command. For example, **passwd** is both a command, as well as a config file. To access the manual for the **passwd** command, one can write:

```
marcel@dell:~$ man passwd
```

or

```
marcel@dell:~$ man 1 passwd
```

However, to access the **passwd** config page manual, one needs to write:

```
marcel@dell:~$ man 5 passwd
```

7.3 **whatis**

As shown in Listing 7.1, each manual page comes with a description section. The **whatis** command searches the man page of the command provided as an argument and displays its description. See the example below:

```
marcel@dell:~$ whatis ls
ls (1)          - list directory contents
```

If the argument can be found in multiple sections, all descriptions are provided in the output, as shown below:

```
marcel@dell:~$ whatis passwd
passwd (1)      - change user password
passwd (5)      - the password file
```

7.4 **apropos**

In English, the word *apropos* means: *concerning, with reference to*. The **apropos** command is called using keywords as arguments. It returns a single line for each man page that contains one or more of the keyword(s) in its **NAME** section.

This command comes handy when you need to search for commands which names you are not certain of. See the example below*:

```
marcel@dell:~$ apropos rename
dpkg-name (1) - rename Debian packages to full package names
mv (1)        - move (rename) files
rename (1)    - renames multiple files
rename (2)    - change the name or location of a file
renameat (2)  - change the name or location of a file
```

**A few extra outputs from this command were omitted for simplicity.*

In this example you can see that the **mv** command can be used to rename a single file, while the **rename** command is normally only used to rename multiple files at once.

7.5 **info**

Man pages were designed in the early seventies. Hence, they can only display simple text and do not take advantage of newer technologies such as hyperlinks. Also, man pages can be quite terse and seldom provide examples.

In order to modernize the system documentation, the GNU project introduced the concept of info pages. Info pages are similar to man pages in which they provide a description of commands, daemons, and config files. However, they differ in two crucial aspects:

1. Info pages provide a much more in depth description of commands and the options they can take. Multiples examples are often provided.
2. Info pages are divided in nodes that can be accessed via hyperlinks*, **To follow a hyperlink, you need to press enter when you cursor is over a node title.* or with the shortcuts provided in Table 7.2.

q	Quits the info page
n	Moves to the next node
p	Moves back to the previous node
u	Go up to the parent node
h	Display a list of shortcuts to navigate info pages (exit it by pressing x)

Table 7.2: Shortcuts to navigate info pages

The choice between using a man page or an info page depends on the user. Some people prefer the terseness of man pages, while other find that it makes it hard to understand. On the other hand, some people like the level of detail present in info pages, as well as the fact that it contain examples, while others find it difficult to find the information they need among multiple nodes.

As a rule of thumb, you may want to read first the man page for a particular command, and only read its info page in case you could not find the information you needed.

EXERCISES

1. Why there is no help page for the **ls** command in bash?
2. Why there is no man page for the **alias** command in bash?
3. What is a built-in command?
4. Why are some commands implemented as built-ins, as opposed to being implemented as binary files?
5. Why are manual pages divided in sections?
6. Explain in which scenarios you would use the **whatis** command.
7. Explain in which scenarios you would use the **apropos** command.
8. what is the difference between a man page and an info page?

9. How can you use the **ls** command to show all files in the working directory (including hidden files), without also showing the implied **.** and **..** directories? *Hint: check its man or info pages.*
10. How can you use the **rm** command to remove an entire directory called **myFolder** in the working directory, while asking for confirmation before deleting each file inside **myFolder**? *Hint: check its man or info pages.*
11. How can you use the **cp** command to make copies of one file called **myFile** in the working directory, saving it into a folder called **myFolder**, only if there are no files called **myFile** inside **myFolder**, or if the file exists, but is older than the new copy you are trying to create. *Hint: check its man or info pages.*

```

1 MKDIR(1) User Commands MKDIR(1)
2
3 NAME
4     mkdir - make directories
5 SYNOPSIS
6     mkdir [OPTION]... DIRECTORY...
7 DESCRIPTION
8     Create the DIRECTORY(ies), if they do not already
9     exist.
10
11     Mandatory arguments to long options are mandatory
12     for short options too.
13
14     -m, --mode=MODE
15         set file mode (as in chmod), not a=rwx - umask
16     -p, --parents
17         no error if existing, make parent
18         directories as needed
19     -v, --verbose
20         print a message for each created directory
21     -Z      set SELinux security context of each created
22     directory to the default type
23     --context[=CTX]
24         like -Z, or if CTX is specified then set the
25         SELinux or SMACK security context to CTX
26     --help  display this help and exit
27     --version  output version information and exit
28 AUTHOR
29     Written by David MacKenzie.
30 REPORTING BUGS
31     GNU coreutils online help: <http://www.gnu.org/
32     software/coreutils/>
33     Report mkdir translation bugs to <http://
34     translationproject.org/team/>
35 COPYRIGHT
36     Copyright 2014 Free Software Foundation, Inc.
37     License GPLv3+: GNU GPL version 3 or later <http://gnu.
38     org/licenses/gpl.html>.
39     This is free software: you are free to change and
40     redistribute it. There is NO WARRANTY, to the extent
41     permitted by law.
42 SEE ALSO    mkdir(2)
43
44     Full documentation at: <http://www.gnu.org/software
45     /coreutils/mkdir> or available locally via: info '(
46     coreutils) mkdir invocation'
47
48 GNU coreutils 8.23 November 2015 MKDIR(1)

```

Listing 7.1: Manual page for the **mkdir** command.

READING AND EDITING TEXT FILES FROM THE SHELL

In Chapter 5, we learned how to create empty files with `touch`, rename them with `mv`, and even delete them using the `rm` command, among other things. However, so far we have not yet covered how to read the information contained in files, nor how to edit them using the terminal.

In this chapter, we will fill this gap by introducing tools that allow us to read text files (**more**, and **less**), as well as to edit text files (**nano**, **vim**).

8.1 READING TEXT FILES

8.1.1 **more**

To read simple text files on your terminal you can use **more**, by calling it with the name of the file you want to read passed as an argument. See the example below:

```
marcel@dell:~$ more poem
Nature's first green is gold,
Her hardest hue to hold.
Her early leafs a flower;
But only so an hour.
Then leaf subsides to leaf.
So Eden sank to grief,
So dawn goes down to day.
Nothing gold can stay.

marcel@dell:~$
```

This tool will display as much of the file as it fits the terminal screen*. To keep reading the file, you need to press the **Space** key. To quite **more** before reaching the end of the file you can press **q**.

more is quite outdated. Its most glaring limitation is that it doesn't allow you to scroll backwards while reading files. You can only scroll forward. So, while we included it here for historical reasons*, even its man page advises users to use the **less** tool instead.

In the provided example, all text from the **poem file was able to fit in the screen.*

some embedded systems still use **more to take advantage of its small size.*

8.1.2 **Less**

The **less** command was created with the main goal of providing backwards scrolling to **more***. It has since become the *de facto* tool to read text files on the terminal. For example, all manual pages are displayed using **less**.

This command is quite different than all other commands we have learned on chapters 5 and 6. So far, issuing commands would normally result in the steps below:

1. The user calls a command which might take options and arguments
2. The command that was called processes the user's input
3. All command's outputs are displayed on the terminal
4. A new shell prompt is made available just below the previous command's output

When calling **less**, on the other hand, the following sequence of steps happens:

1. The user calls **less** providing the name of a text file as an argument, as shown in the example below:

```
marcel@dell:~$ less poem
```

2. **less** starts its own user interface, as shown in Listing 8.1, that takes over the entire terminal screen and displays the beginning of the document
3. The user can use a number of keys to navigate the document
4. The user enters a key to quit **less**' interface
5. A new command prompt is made available below the one in which the user called **less***.

By creating its own user interface, **less** can allow the user to perform actions that would either not be possible or very cumbersome otherwise, such as backwards scrolling and some advanced navigation methods. These actions can be performed using the keyboard. The most important keys to control **less** are shown in Table 8.1*.

Terminal User Interfaces

Many other tools create their own terminal user interface like **less**. Notable examples are the two text editors discussed later

**In fact, its name is a pun on the architecture's minimalist design motto "less is more".*

Note that the contents of the file opened in **less do not remain in the terminal screen after the user quits **less**.*

for a comprehensive list, check **less man or **info** pages.*

```

1 Nature's first green is gold,
2 Her hardest hue to hold.
3 Her early leafs a flower;
4 But only so an hour.
5 Then leaf subsides to leaf.
6 So Eden sank to grief,
7 So dawn goes down to day.
8 Nothing gold can stay.
9
10 poem (END)

```

Listing 8.1: less user interface.

ENTER or ↓	Moves forward by one line
Page Up or SPACE	Moves forward by one screen
y or ↑	Moves backward by one line
Page Down or b	Moves backward by one screen
g	Moves to the beginning of the file
G (Shift + g)	Moves to the end of the file
q	Quits less
h	Help

Table 8.1: Less navigation keys.

in this chapter, **nano** and **vim**, as well as the **top** command covered in Chapter XXX.

You can think of these tools the same way you think of applications such as MS Word, or your browser in a [GUI](#) environment. These tools run on top of your shell, the same way as those applications run on top of your [OS](#). The only difference is that they normally take the whole terminal screen, as opposed to opening in a separate window, and normally can only be controlled via keyboard inputs (as opposed to mouse or touchscreen).

When using tools that create their own interface, it is important to understand that they generally do not understand bash commands. I.e., issuing commands such as **mkdir**, or **ls**, inside these tools will most likely result in an error message.

The ability to search for patterns is also a great feature introduced in **less**. To search for a pattern, you need only to type backslash (\) followed by the desired pattern and press **Enter**. For example, by entering **\folder**, you are shown the first occurrence of the word **folder** starting from the top of the current screen. To keep searching

for the next occurrence of the desired pattern, you can type **n**. To go back to a previous occurrence of the desired pattern, you can type **N** (**Shift + n**).

8.2 EDITING TEXT FILES

In this section we will cover two widely used text editors. **nano** for small edits, and **vim** for larger projects and scripts. There is another widely used text editor in the Linux world called **emacs**, with some really die-hard fans, that is not covered in this book. There are two reasons why **vim** was chosen over **emacs**. First, **vim** is available by default in more distributions. Second, the author itself is a **vim** enthusiast.

8.2.1 *nano*

The simplest way to edit a text file on terminal is by using the **nano** command. To do so, you need only to call **nano** followed by the name of the file you are trying to edit*, such as in the example below:

```
marcel@dell:~$ nano poem
```

In this example, a file called **poem** is opened, resulting in the user interface shown below:

**If no file name is provided, nano will open with a new empty file..*

```

1 GNU nano 2.2.6          File: poem
2
3 Nature's first green is gold,
4 Her hardest hue to hold.
5 Her early leafs a flower;
6 But only so an hour.
7 Then leaf subsides to leaf.
8 So Eden sank to grief,
9 So dawn goes down to day.
10 Nothing gold can stay.
11
12 ^G Get He^O WriteOut^R Read Fi^Y Prev Pag^K Cut Te^C Cur Pos
13 ^X Exit  ^J Justify ^W WhereIs^V Next Pag^U UnCut ^T To Spel

```

Listing 8.2: Nano's user interface.

Once a text file is open in **nano**, you can start editing it using the keyboard. To perform actions such as: saving the file, exiting **nano**, or getting help, you simply need to enter the shortcuts presented at the bottom of the interface*. For example, you can exit **nano** by entering **Ctrl + X**, or you can save the file by entering **Ctrl + O**.

**The ^ symbol stands for the Ctrl button.*

8.2.2 **vi** and **vim**

The **vi*** terminal-based text editor was released for Unix systems more than 40 years ago. It has since being ported to multiple systems and OSs, and many text editors are built upon it. **vi** is, together with text editors derived from it, the most widely used text editor for Linux, and can be found in all Unix and Linux distributions.

**short for visual.*

The most famous text editor derived from **vi**, **vim***, augments the capabilities of **vi** by introducing, among other things:

**short for vi improved.*

- Syntax highlighting for multiple programming languages.
- Spell check in more than 50 languages.
- Multilevel undo and redo. I.e., you can undo and redo multiple edits to the text, as opposed to only the last one.
- More user-friendly interface.

Over time, **vim** has become so ubiquitous in the Linux world that currently it is made available by default in most Linux distributions. In fact, this ubiquitousness has reached the point that, in some distributions, the command **vi** actually calls **vim** instead of **vi**.

In this section, we will cover how to perform basic tasks in **vim**. However, all methods described here also apply to **vi**, unless stated otherwise.

vim interface

Opening **vim** to start editing a file is as simple as opening **nano**. You simply need to call **vim** followed by the name of the file you are trying to edit*. See the example below:

```
marcel@dell:~$ vim poem
```

Not that in Listing 8.3, the **vim** user interface indicates empty lines with tildes (~). Also, **vim** provides the user with important information at the bottom of the terminal screen. Table 8.2 presents a list of all information that is displayed, as well as the corresponding values displayed in Listing 8.3.

If no file name is provided, **vim displays a splash screen with some help information. When you start inserting text, **vim** creates a new file.*

vim modes

Working with **vim** requires understanding its three modes of operation: **command mode**, **insert mode**, and **extended mode**. These modes are described below:

COMMAND MODE In this mode, which is also known as **normal mode**, you can use combinations of keys to enter commands. Commands can be used to copy and paste, delete blocks of text, or undo previous actions, among other things.

```

1 Nature's first green is gold,
2 Her hardest hue to hold.
3 Her early leafs a flower;
4 But only so an hour.
5 Then leaf subsides to leaf.
6 So Eden sank to grief,
7 So dawn goes down to day.
8 Nothing gold can stay.
9 ~
10 ~
11 ~
12 "poem" 9L, 210C                1,1                All

```

Listing 8.3: **vim**'s user interface.

File name	"poem"
Number of lines	9L
Number of characters	210C
Cursor position	1,1 (first line, first column)
Position of the screen	Indicates if you are at the Top , Bottom , or which percentage of text has already been read. In the example, All text is displayed.

Table 8.2: Information displayed in **vim**.

INSERT MODE In this mode you can type new text. Note that this is not the mode you need to be to copy or paste blocks of text.

EXTENDED MODE This mode, which is also known as **last-line mode***, can be used to save the current file, open more files, turn the spell check on and off, and quit **vim**, among other uses.

*The name *extended mode* is due to the fact that it was based on a previous text editor called **ex**.

Each time you start **vim**, it is in **command mode**. To switch to **insert mode**, you can press any of the following keys: **a**, **A**, **i**, **I**, **o**, or **O**. Each key results in starting the **insert mode** at a different position with regards to where the cursor is, as shown in Table 8.3. To switch back from **insert mode** to **command mode**, you need to press **Esc**.

To switch from **command mode** to **extended mode**, you need to press the colon key (:). To switch back, from the **extended mode** to the **command mode**, you need to press **Enter**. In Figure 8.1, you can see how to transition from different modes. Note that it is not possible to switch from **insert mode** to **extended mode**, or vice versa, directly.

i	Starts insert mode at the cursor
I	Starts insert mode at the beginning of the line the cursor is
a	Starts insert mode after the cursor
A	Starts insert mode at the end of the line the cursor is
o	Opens a new line below the cursor and starts insert mode on it
O	Opens a new line above the cursor and starts insert mode on it

Table 8.3: Keys to switch to insert mode.




Figure 8.1: Vi operational modes and the keys required to change modes.

Command mode

In **command mode**, **vim** allows the user to perform a multitude of different operations in a fast and direct way by pressing different keys. In fact, **vim** is so powerful that many experienced programmers prefer to use it even when editors with sophisticated GUIs are available.

Some commands are quite intuitive such as the arrow keys moving the cursor around character by character, or the **Page Up** and **Page Down** keys scrolling the text by one screen at a time. However, some commands can be a bit cryptic, such as **y*** standing for copy. Moreover, some commands require a sequence of keys to be pressed in an specific order to attain the desired outcomes.

In order to help new users to get familiar with **vim**, Table 8.4 presents some of the most used commands for **vim**^{*}. For a comprehensive list of **vim** commands, you can refer to its **man** page.

^{*}It is actually short for yank.

^{*}This table is introduced here to be used as a reference, not to simply be memorized..

Extended Mode

The **extended mode** is mostly used for file operations, such as saving files, opening files, etc. It is also used, however, to configure the behaviour of **vim**. For example, it is in **extended mode** that you can turn spell check on and off. Finally, the **extended mode** is also used to quit **vim**.

On Table 8.5, we show some of the most important operations that can be performed in extended mode. Remember that you need to type colon (:), while in command mode, to reach the extended mode.

u	Undo the latest command. Note that all edits performed each time you enter the insert mode are considered as just one command.
Ctrl+R	Redo changes that were undone using the u command.
dd	Deletes the line the cursor is on. Note that the line is saved on a clipboard which can be later pasted using the p command.
#dd	Where # stands for a number, deletes # lines starting at the one where the cursor is on and saves them on the clipboard.
yy	Copies (yanks) the line the cursor is on to the clipboard.
#yy	Where # stands for a number, copies (yanks) # lines starting at the one where the cursor is on to the clipboard.
v	Turns on visual mode that allows you to specify, using the arrow keys, which part of the text you want to act upon with other commands. For example, you can use v to select a few paragraphs, and then delete them all by pressing d .
p	Pastes the contents of the clipboard after the position the cursor is on.
P	Pastes the contents of the clipboard before the position the cursor is on.

Table 8.4: List of some **command mode** important commands.

:q	Quits vim . If any open files have non-saved edits, an error message is displayed.
:w	Saves the file under its current name. If the file hasn't been given a name yet, an error message appears.
:w file_name	Saves the file as <code>file_name</code> . If <code>file_name</code> is the same names as the currently open file, the file is overwritten
:e file_name	opens a new file, while placing the current file in a buffer. Having multiple files open when the same time can be very convenient. Specially when you need to copy parts of one file, and paste it into another.
:ls	Lists all currently opened files.
:b file_name	Switchs to a file called <code>file_name</code> currently open, i.e., currently in the open buffer. You can use <code>tab</code> to autocomplete the file name. You can also use its buffer number instead of its name.
:b#	Switchs to the previously opened file. This is very convenient when you need to go back and forth between two files.
:set spell spelllang=en_ca	Switchs spellcheck on, and assumes that the current language is Canadian English.
:set nospell	Switchs spellcheck on, and assumes that the current language is Canadian English.

Table 8.5: List of some extended mode important commands.

It is possible to combine commands. for example, you can save and exit **vim** by issuing a **:wq** command, instead of two separate **:w** and **:q** commands..

By default, **vim** tries to prevent users from making mistakes, such as quitting before saving edited files, or overwriting existing files. In case you want to override these security measures, you simply need to add an exclamation point (!) to the end of your command. For example, the command **:q!** quits **vim** without saving any open files.

Reading non simple-text file

All tools presented in this chapter will assume that the files being opened are simple-text encoded using an **ascii** table. I.e., they will translate sequences of 8 bits into characters. For example, a sequence **00100000** is translated to **Space**, **01100001** to

a (lower case), and so it goes. For a complete **ascii** table, see <http://www.asciitable.com/>.

In case you try to open a non-simple text file, such as pdf files or binary applications, using the tools presented in this chapter, you will see a seemingly random sequence of **ascii** characters. This happens because the tool will interpret whichever sequence of bits it finds using an **ascii** table. Even if these bits were not created to represent a simple text.

EXERCISES

1. What is the relationship between the **more** and the **less** commands?
2. How can you search for a particular word in **less**?
3. Cite two advantages of using **vim** over **nano**.
4. Type **vimtutor** on your shell prompt. It will open an iterative tutorial to help you practice your **vim** skills.
5. Which actions can you perform in **vim**'s **extended mode**?
6. Which actions can you perform in **vim**'s **command mode**?
7. Which command you would enter to delete 20 lines of text, starting with the current line?
8. Which command you would enter to be able to select parts of your text using the arrow keys, as opposed to entering the number of lines or words you wish to select?
9. Assume that you have two files in your current working directory called **text1** and **text2**. Explain, including all commands you need to enter, how can you copy some lines from one file, and paste it into another.

Part III

ADVANCED COMMAND LINE INTERFACES

In the following chapters, we discuss Linux file systems, covering both its formatting, as well as its directory hierarchy. Also, we introduce file links, discussing both hard links and soft links, as well as their properties. Following, we introduce file globbing techniques and discuss different wildcards that are frequently used. We also provide a gentle introduction to regular expressions, combined with the **grep** tool. Finally, we discuss important filters and tools, such as **cat**, **sort**, **gawk**, **rename**, and **find**. We finish this part by explaining the important concepts of piping and redirection.

LINUX FILE SYSTEMS

In Linux, the expression **file system** can refer to two distinct concepts, which can lead to some confusion. One of such concepts has to do with the **format** in which data is stored. The other one has to do with the **hierarchy** of directories and subdirectories that are present in a Linux [OS](#). In what follows, we explain these two concepts separately.

9.1 FILE SYSTEM: DATA STORAGE FORMAT

In order to access data in hard-disks, USB devices, optical devices (CD, DVD, etc), etc., the [OS](#) kernel needs to know the format in which this data is written. You can think about this format as the language that has been used to write the data. It specifies how the memory is divided in blocks and how data is stored inside these blocks. Also, it specifies which information about the data, also known as metadata, is stored. Finally, it also specifies in which blocks and what type of metadata needs to be stored.

9.1.1 *Linux File System Evolution*

Originally, the Linux kernel used the MINIX file system*. The first file system format created specifically for the Linux kernel is called the extended file system, or **ext**. Over the years, this file system has been improved upon, leading to three different formats being in use today, as described in what follows:

**MINIX was the Unix-like system Linus Torvalds used as a starting point for creating Linux.*

EXT2 This first update on the **ext** file system works on systems with up to 32 Terabytes (compared to only 2 Gigabytes in ext). Also, it keeps track of different timestamps for when files were last accessed, modified, and changed*. This file system is still widely used for SD cards and USB flash drives.

**I.e., had its metadata, such as its permissions, changed.*

EXT3 This update improves upon **ext2** by introducing a **journalling system**. Under this system, the kernel keeps a journal* of all modifications it needs to make in order to properly save (or delete) data in the memory. In case there is a crash, the system can avoid data corruption by checking the contents of the journal. Journalling is normally avoided in SD cards and USB flash drives because it requires additional memory writes, which decreases their lifespan.

**The journal is saved as a hidden file.*

EXT4 This update improves upon **ext3** by allowing an infinite number of subdirectories, as opposed to “only” 32,000, and by work-

ing with volumes of up to 1 Exabyte. Also, it allows some administrative tasks, such as repairing file systems, a lot faster. Finally, **ext4**, as opposed to **ext3**, allows journalling to be turned on and off.

As pointed above, **ext2** is still widely used for smaller data storage devices. Also, some non-Linux OSs still lack full compatibility with **ext4**, making the use of **ext3** a better approach in some situations. In summary, even though **ext4** is the most up to date system format for Linux systems, there are some scenarios in which using **ext2**, or **ext3** are more appropriate.

It is important to note that other OSs have their own file systems such as **FAT32** and **NTFS** for Windows, as well as **HFS+** and **APFS** for macOS and iOS. Android normally uses **ext4**. A Linux OS normally needs to be installed in a partition formatted with **ext2**, **ext3**, or **ext4**. However, most Linux distributions contain drivers that allow users to work with devices such as SD cards and USB flash drives formatted using other file systems.

9.1.2 Linux File System Format

File systems in Linux are divided in blocks. These blocks can be divided in three major components*: **superblock**, **inode table**, and regular **data**.

SUPERBLOCK In Linux, superblocks store information about the file system, such as: its size, the size of each block, which blocks are full and which blocks are empty, which inodes are taken and which inodes are free, etc.

INODE TABLE The inode table stores information about each folder or file, including its permissions, timestamps (access, modify, change), size, location in memory, etc. Note that the inode table does not store the filename*. Each file is assigned to an inode entry (in other words an inode number) in the inode table. Hence, the OS can quickly retrieve information about any group of files by simply consulting the inode table.

DATA The vast majority of the memory of any device is allocated for the storage of user and system data. To access the data contained in any file, the OS first consults the inode table to check for permissions, as well as for the location of the required data in memory.

**Strictly speaking, there are other components. However, they are omitted here for the sake of simplicity.*

**Folder names and file names are stored in directory files, as discussed in the block Section 9.1.4.*

9.1.3 **stat**

To check the data stored in the inode table for a particular file, the **stat** command can be used. It takes as an argument the name of the

file or folder for which the information is required, as shown in the example below:

```
marcel@dell:~$ stat poem
  File: 'poem'
  Size: 210          Blocks: 8          IO Block: 4096
    regular file
Device: 805h/2053d Inode: 12075464    Links: 1
Access: (0664/-rw-rw-r--)  Uid: (1000/marcel)  Gid: (1000/
    marcel)
Access: 2016-08-01 17:54:19.553690653 -0400
Modify: 2016-07-27 21:49:03.334173611 -0400
Change: 2016-07-27 21:49:03.334173611 -0400
 Birth: -
marcel@dell:~$
```

The **stat** command displays the size of the file in bytes, the number of blocks, the block size, the type of file, a number that identifies the device* in which the file exists (both in hex and in decimal format), its inode number, and how many hard links the file have. Also, the **stat** command displays the access permissions for the file, as well as the name of its user owner and group owner ids. Finally, the **stat** command shows the timestamps for access (file has been opened), modify (contents have been altered), and change (metadata has been altered). Note that it does not display the timestamp for when the file was created, which can be seen by the fact that the **Birth** field is empty. This is a future feature that hasn't yet been implemented.

*Hard-disk, USB stick, DVD, etc..

9.1.4 Directories in Linux

Directories in Linux are nothing else than special files*. They hold tables containing the following information about its files and sub-folders: their names and their corresponding inode numbers. As an example, Table 9.1 below depicts the contents of the directory file for `/home/marcel`.

*As we explain in Section 9.2, everything in Linux is a file.

FILE NAME	INODE
Documents	12583653
Downloads	12583650
Pictures	12583655
Music	12583654
Video	12583652
Seneca.pdf	12583421

Table 9.1: Contents of the `/home/marcel` directory file.

9.1.5 Accessing data and metadata

After learning the format in which Linux file systems are implemented, it is important to understand how the OS performs some basic file operations. In what follows, we present a few examples of file operations, and we describe how they are performed, under the hood, by the OS.

For example, when you issue an `ls` command, the OS needs only to check the names of all entries in the directory file. However, if you issue an `ls -l` command, the OS needs to check the inode table entries for all files in the directory. This is due to the fact that permission access information is stored in inode tables, and not on directory files.

As another example, when a file is deleted, for instance using the `rm` command, the OS simply sets the status of its inode number as free in the superblock.

Retrieving data from files is a bit more complex. In this scenario, the OS needs to navigate its way, starting from the root directory (`/`) until reaching the desired file. See the example below, where a user is trying to access data from a file `/home/marcel/script.sh`

1. The OS starts by checking the inode number 2 in the inode table, which always points to the location of the root directory (`/`).
2. The inode number of `/home` is retrieved from the root directory file.
3. The inode entry of the `/home` directory file is used to retrieve its location in memory.
4. The `/home` directory file provides the inode number for the `/home/marcel` directory file.
5. The inode entry of the `/home/marcel` directory file is used to retrieve its location in memory.
6. The `/home/marcel` directory file is analyzed to retrieve the inode number of `/home/marcel/script.sh`
7. Finally, the system uses the info located in the inode entry of `/home/marcel/script.sh` to retrieve the data in `script.sh`.

This sequence of steps is depicted in Figure 9.1.

9.1.6 Performing actions on files, directories, and inodes

Even though the sequence of steps presented in Figure 9.1 might appear to be more complex than necessary, the Unix file system* was designed with security, speed, and reliability in mind, as the examples below illustrate:

*Which all Unix-like file system, including Linux, follow.




Figure 9.1: Sequence of steps to access a text file.

- When a file moves from one location to another, within the same partition, only the contents of some directory files change. The data itself remains at the same position in memory. Again, this operation is very fast regardless of the size of the files.
- This system keeps metadata, such as access permissions, separated from the data itself. This separation allows, among other things, the data to be stored in an encrypted format, and only decrypted when the user trying to access it has the proper access permissions.
- As stated above, when a file is deleted, the **OS** simply sets the status of its inode number as free. This way, the memory this file used to occupy becomes free for other files to use in the future. Because only an inode entry needs to have its status changed, deleting files is a very fast process regardless of the size of the file.
- This file system design allows file links to be created in a fast and efficient way, as we will discuss in Chapter XXX.

9.1.7 **shred**

When a file is deleted in Linux, using a command such as **rm**, the **OS** simply sets the status of its inode to free, as explained above. However, the data itself still exists in the memory device until it gets rewritten by other data. This can pose a security risk, as there are methods to read data in a memory device even if its inode entry is declared free.

In order to really erase sensitive data from a memory device, you need to replace the bits representing the data with random bits. This can be accomplished with the **shred** command. Calling **shred** with

a file name as an argument, such as in `shred poem`, replaces the bits in the `poem` file with random bits. If the option `-u` is given, such as in `shred -u poem` the `shred` command first replaces the bits from the `poem` file with random ones, and then deletes the file.

9.2 DIRECTORIES HIERARCHY

A crucial aspect of the Linux file system is that it implements just about everything as a file* In other words, a text file is a file, an application binary is a file, a directory is a file, processes are managed as files, and even pieces of hardware such as a printer or a mouse are represented by files.

**This approach, first proposed for Unix, is employed in all Unix-like OSs.*

The idea of representing everything as a file greatly simplifies the design of the OS. In fact, following this methodology, the role of the OS is to simply pass text and bytes back and forth between different files. For example, to print a text file using a printer, the OS must pass the information in the text file to the file that represents the printer device.

A clear requirement for an OS to function properly, given this “everything is a file” design methodology, is that it needs to know where all the files necessary for its proper operation are. In order to accomplish this goal, the Linux File system Hierarchy Standard (FHS) established the basic directory hierarchy that Linux distributions should use*. Having all distributions using the same directory hierarchy layout results in a series of advantages, among which it is worth to cite:

**Most distributions deviate slightly from the proper standards. To see your distribution basic directory hierarchy, enter `man hier`.*

- It helps different systems to communicate effectively with one another.
- It allows applications to be designed for multiple distributions without requiring any modifications.
- System administrators can easily find system configuration and log files when dealing with multiple distributions.

In what follows, on Table 9.2, we present a list of the most important basic directories in a Linux file system hierarchy, as specified by the FHS standard. We also explain what types of files these directories should hold.

Directory	Contents
/	The root directory. It gets its name because it is from it that the whole directory tree starts. Note that even USB sticks, external hard disks, and optical devices such as DVDs, are mounted on top of the root directory. This is an approach very different than what is done on Windows, where each memory device is mounted as a one letter file system (C:, D:, E:, etc).

/bin	Contains executable programs (binaries) which are needed in single user mode and to bring the system up or repair it. Many important commands, such as ls , mv , and even the bash shell itself are implemented as binary files in this directory.
/boot	Contains static files for the boot loader. It holds only the files which are needed during the boot process.
/dev	Holds files representing physical devices such as a printer, a network card, etc.
/etc	This directory is a place for files that did not fit into other previously defined directories. Its name is derived from the latin expression <i>et cetera</i> . It mostly contains system configuration files.
/home	Holds all user's home directories. For example, the home directory of a user called john will likely be located in /home/john .
/lib	Contains shared libraries that are necessary to boot the system. Shared libraries are compiled pieces of code that can be used by multiple programs in order to achieve a well-defined goal.
/media	Contains mount points for removable media such as CDs, DVDs, SD cards, or USB sticks. I.e., by default these memory devices will be mounted as subfolders inside this directory.
/mnt	Countains mount points for temporarily mounted devices. This directory is the point where a sysadmin should use to manually mount devices.
/opt	Holds program files for software installed without using the apt-get or yum tools (covered on Chapter XX). In other words, software installed without using repositories.
/proc	holds files with information regarding proccess. It contains runtime system information such as: system memory, hardware configuration, etc.
/root	Home directory for the root user (sysadmin)
/sbin	Holds executables (binaries) for basic commands that are normally executed by the OS or by users with root access, such as ifconfig , fdisk , etc.
/srv	Holds files containing data that is served by the system, as opposed to directly accessed. It is where files served via ftp , cvs , or even http , should go, for example.

/tmp	Holds temporary files. Temporary files are normally created and deleted without the user's input. They can be used to store information necessary during boot time, to retain important information when some types of software are open, or for providing users with recovery file options after crashes.
/usr	Contains a number of subfolders such as /bin , /sbin , and even /lib . These folders will host binaries and shared libraries deemed non-essential. It also holds binaries and libraries for software installed through the repositories.
/var	This directory contains files which may change in size, such as spool and log files. Some systems use this folder to hold files whose size can vary, containing data that are served by the system. Other systems, use the /srv folder.

Table 9.2: Basic directory contents according to the [FHS](#).

EXERCISES

1. Are there good reasons to use the **etx2** or **etx3** formats, as opposed to the newest **etx4** format? If so, provide a few.
2. With regards to file systems, what is a journalling system?
3. Can a Linux [OS](#) read memory devices formatted with a FT32 or an NTFS file system?
4. Where in the file system is the information about which inodes are free, and with inodes are taken? In the superblock, in the inode table, in the directory files, or in the data section of the memory device?
5. Where are the names of the files stored? In the superblock, in the inode table, in the directory files, or in the data section of the memory device?
6. What is the difference between the **Modify** and **Change** timestamps?
7. What type of information is stored in directory files?
8. It is said that Linux systems follow a “everything is a file” design. Explain what does this design methodology mean.
9. What happens with the data in a particular file when the file is deleted from the system using a **rm** command.

10. Why is it important for different Linux distributions to follow a standardized directory hierarchy?
11. What is the difference between the **/bin** and **/sbin** directories?
12. What is the difference between the **/media** and **/mnt** directories?
13. What type of information is stored in the **/lib** directory?
14. Where are temporary files stored in a Linux distribution that follows the [FHS](#) standard?
15. Assume you have a web server providing/collecting information to/from users using an **http** protocol. Where would you store the contents you are providing to your users?

FILE LINKS

Hyperlinks are one of the most common features of pages in the world wide web. As you use your browser* to navigate your way to whatever it is you are trying to find, you click on hyperlinks to go from one page to another

**Edge, Safari, Chrome, Firefox, etc.*

File links are not that different. The purpose of a file link is to direct you to another file. Unix-like systems such as Linux have two different types of links: hard links, which were the original way of creating links in Unix, and soft links, also known as symbolic links, which were created to overcome some limitations of hard links*. In what follows, we cover both types of links.

**Soft links are similar to shortcuts in Windows.*

10.1 HARD LINKS

Following our discussion in Section 9.1, it is clear that when a file is created, the following steps are taken by the OS:

- A new entry is added to its parent directory file containing its name and inode number.
- The inode table is updated with the proper information.
- The file data is stored in a particular location of the memory device pointed by the inode table.

When a hard link is created, the OS simply adds a new entry to the chosen directory file with the same inode number as the original file. Also, this procedure will automatically increment the hard link count associated with the file by one.

In Figure 10.1, you can see a diagram showing how a file system is affected by the addition of a hard link. In this example, a file called **tests** initially exists inside a folder **/OPS105**, as it can be seen in 10.1a. Then, a hard link to this file, called **tests_ln**, is created inside a folder **/SRT311**. You can see in 10.1b how both the original file **tests** and the hard link **tests_ln** point towards the same inode entry.

10.1.1 **ln**

To create a hard link, you simply need to use the **ln*** command, followed by two arguments: the name of the file you are creating a link for, including its absolute path, and the name of the link, including the path to the folder in which you want to place it. See the example in page 65.

**Short for link.*



(a) File system before the hard link `tests_ln` was created.(b) File system after the hard link `tests_ln` was created.

Figure 10.1: Schematics diagram of a file system before and after the addition of a hard link.


```

1 marcel@dell:/$: tree
2 .
3 |-- OPS105
4     |-- students
5     |-- tests
6 |-- SRT311
7     |-- students
8 2 directories, 3 files
9 marcel@dell:/$ln /OPS105/tests /SRT311/tests_ln
10 marcel@dell:/$tree
11 .
12 |-- OPS105
13     |-- students
14     |-- tests
15 |-- SRT311
16     |-- students
17     |-- tests_ln
18 2 directories, 4 files
19 marcel@dell:/$ stat OPS105/tests
20   File: 'OPS105/tests'
21   Size: 128          Blocks: 4          IO Block: 4096
22   regular empty file
23 Device: 806h/2054d Inode: 14680422   Links: 2
24 Access:(0664/-rw-rw-r--) Uid:(1000/marcel) Gid:(1000/marcel)
25 Modify: 2016-08-13 14:57:16.437570016 -0400
26 Change: 2016-08-13 14:59:33.976563502 -0400
27 Birth: -
28 marcel@dell:/$ stat SRT311/tests_ln
29   File: 'SRT311/tests_ln'
30   Size: 128          Blocks: 4          IO Block: 4096
31   regular empty file
32 Device: 806h/2054d Inode: 14680422   Links: 2
33 Access:(0664/-rw-rw-r--) Uid:(1000/marcel) Gid:(1000/marcel)
34 Modify: 2016-08-13 14:57:16.437570016 -0400
35 Change: 2016-08-13 14:59:33.976563502 -0400
36 Birth: -
37 marcel@dell:/$

```

Note that the file **tests** first appears on line 5, under the **OPS105** folder. After the **ln** command, a hard link **tests_ln** appears in the **SRT311** folder (see line 17). Moreover, by using the **stat** command, we can verify that both files have the same inode number and a links count of 2 (see lines 22 and 31).

10.1.2 *Hard Link Properties*

A hard link is indistinguishable from the original file itself. I.e., after the creation of hard links, the OS cannot tell which one was the original file and which ones are its hard links. They are all treated the same way.

An interesting, and somewhat counterintuitive, feature of hard links is that they would allow you to access the contents of the original file, even after the original file is deleted. This is easy to understand by looking at Figure 10.1b. In it, you can see that the hard link provides you with access to the original file's inode entry even if the original file gets removed. In fact, the inode entry for any file is considered taken until its hard link count goes to zero. A consequence of this behaviour is that, in order to successfully delete a file, you need to delete the original file, as well as all of its hard links.

10.1.3 *Hard link limitations*

Hard links have two important limitations:

- A hard link cannot reference a file outside its own file system. This means that a link cannot reference a file that is not on the same disk partition as the link itself. The reason behind this limitation is that each partition has its own inode table. Hence, providing an inode number is not enough for the OS to resolve in which partition is the file located.
- A hard link may not reference a directory. This limitation prevents users from accidentally creating link loops while retrieving files. Imagine a scenario where a file `/var/log/syslog` exists. Now, let the user create a hard link for the `/var` directory using: `ln /var /var/log`. In this scenario, the directory `/var` would be both a parent, and a child of `/var/log`. Hence, the system would find the files: `/var/log/syslog`, `/var/log/var/log/syslog`, `/var/log/var/log/var/log/syslog`, and so it goes.

10.1.4 *Hard Link usage*

Hard links are not as common as soft links, but they do appear in two distinct scenarios:

- To save space while performing automatic backups. Whenever a file to be saved as a backup hasn't changed, you can save memory by simply creating a hard link pointing to the last backup that was modified, as opposed to creating a full copy of it. For example, suppose that your system creates a backup version of a file `OPS105.zip` every month. Also, suppose

it hasn't changed since April 2016. Using hard links, you can create links called **OPS105_04_2016.zip**, **OPS105_05_2016.zip**, **OPS105_06_2016.zip**, etc. Without having to waste any memory. Also, you can delete any one of the hard links, and even the original file, without ending up with a broken link*.

**As long as at least one hard link remains.*

- To allow different commands to call the same executable. For example, in Ubuntu, **bzcat**, **bunzip2**, and **bzip2**, are hard links to the same compressing tool.

10.2 SOFT LINKS

Soft links, also known as symbolic links, or even as symlinks, were created to overcome the limitations of hard links. They work by creating a special type of file that contains a text pointer to the referenced file or directory. In this regard, they operate in much the same way as a Windows shortcut*.

**They predate Windows shortcuts by many years, though.*

In Figure 10.2, you can see a diagram showing how a file system is affected by the addition of a soft link. In this example, which is very similar to the previous one, a file called **tests** initially exists inside a folder **/OPS105**, as it can be seen in 10.2a. Then, a soft link to this file, called **tests_ln**, is created inside a folder **/SRT311**. You can see in 10.2b that the soft link points to a new inode. Also, the file pointed by this new inode, in its turn, points to the address of the original file **/OPS105/tests**.

10.2.1 **ln -s**

Creating a soft link is very similar to creating a hard link. You simply need to use the **ln** command with the **-s** option, followed by two arguments: the name of the file you are creating a link for, including its absolute path, and the name of the link, including the path to the folder in which you want to place it. See the example below (which continues in Page 69):

```
1 marcel@dell:/$: tree
2 .
3 |-- OPS105
4     |-- students
5     |-- tests
6 |-- SRT311
7     |-- students
8 2 directories, 3 files
9 marcel@dell:/$ln -s /OPS105/tests /SRT311/tests_ln
10 marcel@dell:/$tree
11 .
12 |-- OPS105
```

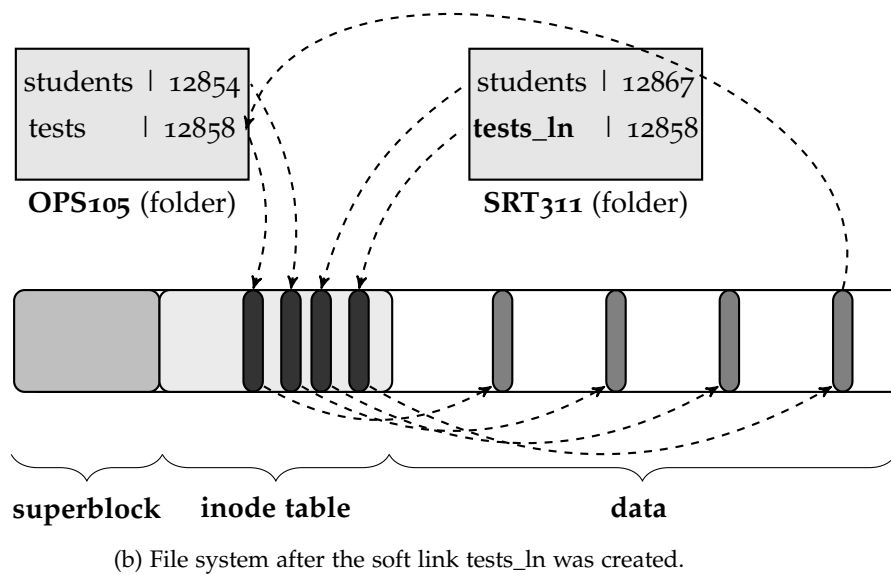



Figure 10.2: Schematics diagram of a file system before and after the addition of a soft link.

```

13 |-- students
14 |-- tests
15 |-- SRT311
16 |-- students
17 |-- tests_ln -> /OPS105/tests
18 2 directories, 4 files
19 marcel@dell:/$ stat OPS105/tests
20 File: 'OPS105/tests'
21 Size: 128          Blocks: 4          IO Block: 4096
   regular empty file
22 Device: 806h/2054d Inode: 14680422   Links: 1
23 Access:(0664/-rw-rw-r--) Uid:(1000/marcel) Gid:(1000/marcel)
24 Access: 2016-08-13 14:57:16.437570016 -0400
25 Modify: 2016-08-13 14:57:16.437570016 -0400
26 Change: 2016-08-13 14:59:33.976563502 -0400
27 Birth: -
28 marcel@dell:/$ stat SRT311/tests_ln
29 File: 'SRT311/tests_ln'
30 Size: 128          Blocks: 4          IO Block: 4096
   regular empty file
31 Device: 806h/2054d Inode: 14680425   Links: 1
32 Access:(0664/-rw-rw-r--) Uid:(1000/marcel) Gid:(1000/marcel)
33 Access: 2016-08-13 14:57:16.437570016 -0400
34 Modify: 2016-08-13 14:57:16.437570016 -0400
35 Change: 2016-08-13 14:59:33.976563502 -0400
36 Birth: -
37 marcel@dell:/$

```

Note that the file **tests** first appears on line 5, under the **OPS105** folder. After the **ln** command, a soft link **tests_ln** appears in the **SRT311** folder (see line 17) pointing towards **/OPS105/tests**. Moreover, by using the **stat** command, we can verify that these files have different inode numbers, and that their links count is only 1 (see lines 22 and 31).

Absolute vs Relative paths for links

It is possible to use a relative path, as opposed to an absolute path, when creating a soft link. However, this practice is not recommended.

The reason for absolute paths being preferred when creating soft links is that the link must contain a path that can lead from it to the original file. I.e., it must be given a relative path with regards to the location of the link, not with regards to the location of the current working directory at the time the link was created. This can be quite confusing and lead to errors.

Another reason for preferring absolute paths, over relative paths, is that soft links created with relative paths become broken when the link is moved to a different directory. Soft links created with an absolute path, on the other hand, still work after moving to different locations.

Note that for hard links, it doesn't matter if you create them with a relative path, or with an absolute path. Once created, they will keep pointing towards the proper inode even if they are moved.

10.2.2 Soft link properties

When you try to access the contents of a soft link, the OS redirects you to the original file. For example, if you try to open a soft link in **vim**, the original file is opened. If any edits are made to it, the original file changes*.

If the original file is deleted before the symbolic link, the link will continue to exist, but will point to nothing, resulting in a broken link. In many Linux distributions, the **ls** command displays broken links in a different color, such as red, to indicate that they are broken.

Note that, commands such as **mv, **cp**, or **rm** will affect the link itself, not the original file.*

EXERCISES

1. Can you access a file using a soft link, after the soft link is moved to a different folder after being created?

Given the directory tree shown in Figure 10.3, answer the following questions:




Figure 10.3: Directory tree for questions 2 and 3.

2. Create a hard link for the file **ops105.sh**, inside the folder **Seneca**. Assume that your current working directory is the root (/).

3. Create a soft link for the file **srt311.py**, inside the folder **Seneca**. Assume that your current working directory is the root (/).
4. Can you create a hard link to another hard link? How about creating a soft link that points towards another soft link?
5. What happens to a soft link when the original file it points to is deleted?
6. What happens to a hard link when the original file it points to is deleted?
7. In which scenarios are hard links preferred over soft links? In which scenarios are soft links preferred over hard links?
8. why is it important to use absolute paths while creating soft links?

FILE GLOBBING: USING WILDCARDS

System administrators, very often, have to act on multiple files at once. For example, they might want to move all **.pdf** files from one folder to another. As another example, they might want to list all files within a particular directory that have the string* **net** as part of its file name, such as **netcat**, **netstat**, or **networkctl**, in the **/bin** directory.

**A string is a sequence of characters.*

Up until now, we have only learned how to run commands with specific file names passed as arguments. For example, **mv grades.xls old_grades.xls** renames a file called **grades.xls** in the current working directory. As another example, **rm midterm.pdf final.pdf** deletes both files **midterm.pdf** and **final.pdf**.

This approach used in the previous paragraph, i.e., specifying the name of each file we want to work upon in the arguments list, works well when we are dealing with a small number of files. However, it is clearly impractical in scenarios where we need to deal with dozens, hundreds, or even thousands of files at once. For these scenarios, we can use wildcards* that can be used to match all files that follow a particular pattern.

**Wildcards get their name because they work as a Joker card in some card games. They can represent different things, depending on the situation.*

In the **Bash** shell, there are three wildcards that can be used for different purposes:

STAR - * : This wildcard can replace any number of characters from file names, including none.

QUESTION MARK - ? : This wildcard can replace one, and only one, character from a file name.

SQUARE BRACKETS - [] : Square brackets can replace one, and only one, character from a file name. However, the space occupied by the square brackets can only be occupied by characters from a specified list.

In what follows, we cover each one of these three wildcards and provide a number of examples.

11.1 STAR - *

The star wildcard is by far the most used one. One particularly common usage of this wildcard is by itself, when it expands to all files in the current working directory. For example, by issuing a **rm *** command, all files in the working directory are deleted.

Another very common use of this wildcard is to allow users to specify all files with a specific file name ending. See the following example:

```
marcel@dell:~$ tree
.
|-- filea.pdf
|-- fileb.pdf
|-- pdf_folder
1 directory, 2 files
marcel@dell:~$ mv *.pdf pdf_folder
marcel@dell:~$ tree
.
|-- pdf_folder
    |-- filea.pdf
    |-- fileb.pdf
1 directory, 2 files
```

In this example, two **.pdf** files in the working directory are sent to a folder called **pdf_folder**. Note that the files are not specified by name, but rather by a file globbing expression ***.pdf**. If there were more **.pdf** files in the working directory, they would also be transferred to the **pdf_folder**.

The star wildcard can also be used to specify all files starting with a particular string, or even to specify files that start with a particular string and end with another string. See the examples below.

```
marcel@dell:~$ tree
.
|-- script1.sh
|-- script2.sh
|-- scripts <-- directory
|-- script_test
|-- test1
|-- test2
1 directory, 5 files
marcel@dell:~$ rm test*
marcel@dell:~$ tree
.
|-- script1.sh
|-- script2.sh
|-- scripts <-- directory
|-- script_test
1 directory, 3 files
marcel@dell:~$ mv script*.sh scripts
marcel@dell:~$ tree
.
|-- scripts
|   |-- script1.sh
```

```
| |-- script2.sh
|-- script_test
1 directory, 3 files
```

In this example, all files that start with the string **test** are deleted at once using the command **rm test***. Also, all files that start with the string **script**, and end with **.sh**, such as **script1.sh**, **script2.sh**, etc., are moved to a folder called **scripts**, using a command **mv script*.sh scripts**.

File Globbing: Under the Hood

The file globbing mechanism is quite interesting. For instance, you may have assumed that the commands **rm** and **mv**, when we entered **rm test*** and **mv script*.sh** above, received **test*** and **script*.sh** as arguments, respectively. However, this is incorrect. In fact, it was the **Bash** shell that expanded these file globbing expressions into all files that match them, before calling the **rm** and **mv** commands.

In other words, the shell expanded these expressions so that, from the point of view of the **rm** and **mv** commands, they were called with: **rm test1 test2** and **mv script1.sh script2.sh**. This is clearly a smart system design choice, as it requires only the shell to implement a file globbing algorithm, as opposed to requiring all commands to implement it individually.

11.2 QUESTION MARK - ?

The question mark, as we said before, replaces one character, and one character only. It is frequently used to match sequences of file names that end on numbers, as seen in the following example:

```
marcel@dell:~$ ls
script1 script2 script2a script_test
marcel@dell:~$ rm script?
marcel@dell:~$ ls
script2a script_test
```

In this example, the files **script1** and **script2**, were deleted from the current working directory. The files **script2a** and **script_test**, on the other hand were left untouched.

Note that you can use more than one wildcard at once, for example, if we had used **rm script??** in the example above, we would have deleted the file **script2a**, while all other files would be left untouched.

11.3 SQUARE BRACKETS - []

Square brackets can be used to substitute one character, and one character only, by a set of characters or numbers, or even by a range of characters or numbers. Note that they are similar to the question mark wildcard in which they substitute one, and only one character. However, they are more restrictive, as they enforce a list of options for possible replacement. See the following example:

```
marcel@dell:~$ ls
script1 script2 scripta scriptb
marcel@dell:~$ rm script[1234]
marcel@dell:~$ ls
scripta scriptb
```

In this example, the files **script1** and **script2**, were deleted from the current working directory, while **scripta** and **scriptb**, were not. Note that if there were files with names **script3** or **script4**, they would also be deleted from the current working directory.

The expression **[1-4]** could have been used in the example above to replace **[1234]**. It means: any number in the range between 1 and 4. We could also have used **[a-z]** or **[A-Z]** to replace any lowercase or uppercase alphabetical character, respectively*. It is possible to select a list of forbidden characters, as opposed to a list of acceptable characters. To do so, you simply need to start the list with an exclamation mark (!). For example, to allow anything but lowercase characters to replace a particular character in a file name, you can use **[!a-z]**.

In order to replace any alphabetical character, regardless of its case, we could have used the expression **[a-zA-Z], or **[A-z]**.*

11.4 ESCAPING SPECIAL CHARACTERS

The shell will, by default, assume that the characters *****, **?**, and square brackets **[]** are to be expanded as wildcards. However, there might be scenarios in which your file names have these characters. In order to match these specific characters, as opposed to use them as wildcards, you simply need to escape them with a backslash character *****. See the example below:

```
marcel@dell:~$ ls
script* script1 script2 scripta scriptb
marcel@dell:~$ rm script\*
marcel@dell:~$ ls
script1 script2 scripta scriptb
```

In this example, only the file **script*** was deleted with the command **rm script***. Note that if we had used **rm script***, instead, all files would have been deleted from this directory.

The backslash character also needs to be escaped. For example, to create a file called **script, we should enter **touch script**.*

EXERCISES

For the exercises below, create an empty directory and run the following command within this directory: **touch myscript script script.sh script1.sh script2.sh scripta.sh scriptb.sh mytest test test.sh testa.sh testb.sh**. Also, create subfolders called **scripts** and **tests** with a **mkdir scripts tests** command.

For each question, you need to issue a single command that will perform the required task. The use of wildcards is required. After each question, make sure to restore the directory back to its initial state by running **rm ***, and then the **touch** and **mkdir** commands specified above.

EXERCISES

1. Delete all files containing the string **script**.
2. Move all files ending in **.sh** to the **scripts** folder.
3. Delete the files **script1.sh script2.sh**.
4. Move the files **scripta.sh**, **scriptb.sh**, and **script1.sh** to the **scripts** folder.
5. Move all files containing the string **test** to the **tests** folder.
6. Delete all files containing the string **test** in it that end with **.sh**.
7. Delete all files from this directory.

GREP AND REGULAR EXPRESSIONS

I have, many times in the past, googled excerpts of song lyrics that got stuck in my head, in the hope that I would find the rest of it, or at least the song's title. Similarly, system administrators can find themselves looking for a specific file whose name they cannot recall, but they can remember excerpts of its contents. In other cases, the sysadmin might want to find in which line of a given file (which can be pretty long), a particular string occurs*.

To help with the scenarios presented in the previous paragraph, most Linux distributions come by default with a tool called **grep*** to search for specific strings, or string patterns, inside text files.

In what follows, we will cover some basic usage of the **grep** tool. Following, we will introduce the concept of regular expressions, which are often used while performing grep searches.

**This is a very common scenario when dealing with configuration files.*

**Short for global regular expression print.*

12.1 **grep**

The **grep** command is called with the following syntax:

```
grep "STRING" LIST OF FILENAMES
```

I.e., it takes at least two arguments: first, the string of characters we are looking for*, and then the names of the files in which we are performing this search on. In its output, **grep** prints on the terminal the lines in the specified files in which the string was found. See the example below:

**It is a good practice to always write the string of characters inside double quotation marks.*

```
marcel@dell:~$ more poem
Nature's first green is gold,
Her hardest hue to hold.
Her early leafs a flower;
But only so an hour.
Then leaf subsides to leaf.
So Eden sank to grief,
So dawn goes down to day.
Nothing gold can stay.
marcel@dell:~$ grep "gold" poem
Nature's first green is gold,
Nothing gold can stay.
marcel@dell:~$
```

Note that we can use file globbing expressions, as opposed to writing a complete list of files. For example, if we wanted to search for lines

in all files in the current working directory containing the string **gold**, we could have used **grep "gold" ***.

12.1.1 *grep options*

The **grep** command offers to system administrators a plethora of options. For instance, it can print line numbers together with the lines in which the required string was found. As yet another example, it can print just the name of the file in which a particular string was found, which is particularly useful when using file globbing techniques. A small list of **grep** options is listed below in Table 12.1*.

For a complete list, check **grep's manual page.*

-i	grep is, by default, case insensitive. Using this option makes it become case insensitive.
-w	Searches for isolated words. Without this option, a grep search for a string "cat" would also return lines with words such as catterpillar , catch , concatenate , etc. Using this option, grep looks for strings that are complete words on their own right.
-v	Reverse grep. Displays only the lines in which there are no matches.
-l	Displays only the name of the files in which at least one match was found.
-n	Shows the line numbers in front of the lines in which matches were found.

Table 12.1: **grep** options.

12.2 REGULAR EXPRESSIONS

Sometimes we want to search for strings that follow a particular pattern, instead of specific strings. For example, postal codes in Canada always follow the pattern **LNL NLN**, where each **L** stands for an uppercase alphabetical character, and each **N** stands for a number between 0 and 9. So, **M3J 3M6**, **M5V 1J1**, and **T6G 2R3** are valid postal codes, whereas **ABC 123**, **A12345**, and **456-1234** are not.

Clearly, there are more scenarios in which we might want to search for strings that follow a particular pattern. Among these: phone numbers, addresses, social insurance numbers, emails, are just a handful of examples.

In order to work with patterns, as opposed to specific strings, **grep** makes use of regular expressions, also known as **regex**. A **regex** normally contains a number of literal characters, as well as special characters, that are interpreted in order to delineate a desired pattern. In

Table 12.2, we introduce a few [regex](#) special characters. Following, we present a few examples.

<code>\</code>	If applied to an special character, it indicates that the special characer should be treated as a regular character. <i>ex: <code>"*" matches the star character.</code></i>
<code>\s</code>	Matches a space.
<code>^</code>	Forces the pattern to match at the beginning of the line. <i>ex: <code>"^A"</code> provides no matches in a line: <code>"an A"</code>.</i>
<code>\$</code>	Forces the pattern to match at the end of the line. <i>ex: <code>"\$\" provides no matches in a line: <code>"No periods here"</code></code></i>
<code>*</code>	Matches the preceding character 0 or more times. <i>ex: <code>"bo*" provides matches in lines with <code>"b"</code>, <code>"bo"</code>, <code>"boo"</code>, etc.</code></i>
<code>\+</code>	Matches the preceding character 1 or more times. <i>ex: <code>"bo\+" provides matches in lines with <code>"bo"</code>, <code>"boo"</code>, <code>"booo"</code>, etc., but not with <code>"b"</code>s not followed by at least one <code>"o"</code></code></i>
<code>\?</code>	Matches the preceding character 0 or 1 time. <i>ex: <code>"bo\?" provides matches in lines with <code>"b"</code> or <code>"bo"</code>. Note that, unless the option <code>-w</code> is used, <code>"boo"</code> would also provide a match, as it contains the string <code>"bo"</code></code></i>
<code>\{N\}</code>	Matches the preceding character exactly N times. <i>ex: <code>"bo{2}" provides matches in lines with <code>"boo"</code>, but not lines with the character <code>"b"</code> by itself or with <code>"booo"</code></code></i>
<code>\{N,M\}</code>	Matches the preceding character N to M times.
<code>\{N,\}</code>	Matches the preceding character N or more times.
<code>.</code>	The period character matches any single character. <i>ex: <code>"c.t"</code> provides matches in lines with <code>"cat"</code>, <code>"cut"</code>, etc.</i>
<code>[]</code>	Square brackets matches single characters, as long as they belong to a specified list. <i>ex: <code>"[cae]t"</code> provides matches in lines with <code>"cat"</code> and <code>"cet"</code>, but not with <code>"cit"</code>, <code>"cut"</code>, etc.</i>

Table 12.2: Special characters for [regex](#).

Note that these special characters are often employed together. For example, the [regex](#) `"[0-9]\{5\}"` matches any number with 5 digits. As another example, the [regex](#) `".*"` matches everything.

12.2.1 Examples of **grep** commands using [regex](#)

Email addresses always have the following structure: they start with a few characters, followed by an `@`, followed by a few more characters, followed by a period sign, an finished by a few more characters. For example, `john@mycomman.com`, `jane@mycollege.ca`. To capture

lines containing emails from a file called **students**, we can use the following command:

```
grep "[a-z]*@[a-z]*\.[a-z]*" students
```

Note that in this example, emails containing numbers, underlines, or period signs before the @ character would not be detected. In order to detect these emails, we could change our command to to:

```
grep "[a-z0-9\._]*@[a-z0-9\._]*\.[a-z]*" students
```

Postal codes in Canada, as discussed at the beginning of this section, follow the pattern **LNL NLN**. Hence, a simple **grep** command that would capture lines containing postal codes in the **students** file would be:

```
grep "[A-Z][0-9][A-Z] [0-9][A-Z][0-9]" students
```

Regular expressions facts

- In some ways, regular expressions are similar to file globbing expressions. However, they tackle a very different problem. Indeed, while file globbing deals with matching complete file names, regular expressions aim at finding strings within large chunks of text. As a result, they are implemented very differently.
- Regular expressions are used in many areas. In fact, most search and find tools implement them. Also, many programming languages such as Perl, Python, and JavaScript, among others work with regular expressions.
- Regular expressions can be quite complex. In fact, a number of books have been dedicated solely to them. In our case, we will simply cover the basics.

EXERCISES

For all questions below, you should write **grep** commands to capture all lines containing the required patterns without capturing any lines not containing the required patterns. Use the following **students** text file (copy and paste it into a text file):

```
marcel@dell:~$ more students
John Smith
(416) 123-2345
M4Z 2P3
john.smith@myseneca.ca
K23a!5

Mohammad Ali
(905) 231-3381
N3P 4A1
mali@myseneca.ca
jE_3sa4G

Akira Kurosawa
(905) 872-1221
M4S 1X4
akira@myseneca.ca
2S!aTe6

Priyanka Singh
(416) 431-3231
M3S 5N2
psingh4@myseneca.ca
X32Dsa
```

1. Student names. Doesn't contain special characters, just alphabetical characters and spaces.
2. Telephone numbers. (NNN) NNN-NNNN
3. Postal codes. LNL NLN
4. Student emails. Note that all students should have an email ending with @myseneca.ca
5. Passwords. Only have alphanumerical characters, as well as the special characters _, !, and ?. Should have between 6 to 8 digits.

BASIC FILTERS

In this chapter, a number of filters that can be used on text files is presented. These filters can be used to join (concatenate) multiple text files (**cat**), show only the first few lines of a text file (**head**), show only the last lines of a text file (**tail**), sort the contents of a text file based on some criteria pertaining one of its columns (**sort**), remove unwanted columns of a text file (**cut**) or even make small changes to a text file such as substituting one word by another (**sed**).

13.1 **cat**

The **cat*** command can be used for two distinct goals:

**Short for concatenate.*

- To display the contents of a short text file
- To concatenate the contents of multiple text files

For the first goal, **cat** works in a very similar way to **more**. You need only to call it together with the name of the file you wish to display. The contents of the file are then displayed on the terminal and you are granted with a new command prompt. In this regard, the major difference between **cat** and **more** is that, for long text files, **cat** doesn't provide the user with a way to navigate the file page by page. The whole file is displayed at once. See the example below:

```
marcel@dell:~$cat poem
Nature's first green is gold,
Her hardest hue to hold.
Her early leafs a flower;
But only so an hour.
Then leaf subsides to leaf.
So Eden sank to grief,
So dawn goes down to day.
Nothing gold can stay.
marcel@dell:~$
```

For the second goal, you need only to call **cat** together with a list of files you are trying to concatenate. The shell will display the contents of all the required files on your terminal one on top of another, and then providing you with a new command prompt. See the following example, where the following text files were used:

```
Nature's first green is gold,
Her hardest hue to hold.
Her early leafs a flower;
But only so an hour.
Then leaf subsides to leaf.
So Eden sank to grief,
So dawn goes down to day.
Nothing gold can stay.
```

Listing 13.1: Poem1 text file

```
What's in a name? That which we call a rose
By any other name would smell as sweet.
```

Listing 13.2: Poem2 text file

```
marcel@dell:~$ cat poem1 poem2
Nature's first green is gold,
Her hardest hue to hold.
Her early leafs a flower;
But only so an hour.
Then leaf subsides to leaf.
So Eden sank to grief,
So dawn goes down to day.
Nothing gold can stay.
What's in a name? That which we call a rose
By any other name would smell as sweet.
marcel@dell:~$
```

It is important to note that the original files are not modified in any way. The shell simply displays the concatenation of their contents on the terminal, without altering them. We will cover methods to create text files based on the results of tools such as **cat** in Chapter XXX.

13.2 head

Sometimes, a system administrator or developer is not interested in all the contents of a text file, but rather in just a part of it. For example, most professional source codes start with a few lines describing its functionality. For such scenarios, opening the entire file is simply a waste of time.

A tool called **head** is provided for cases in which only the first few lines of a text file is of interest. This tool behaves nearly identically to the first usage of the **cat** described previously. The only difference is that it only displays, by default, the first ten (10) lines of the file.

In case you want to display a number of lines different than ten, this command can be called with the **-n** option, followed by the number of lines you wish to display. See the example below:

```
marcel@dell:~$ head -n 4 poem1
Nature's first green is gold,
Her hardest hue to hold.
Her early leafs a flower;
But only so an hour.
marcel@dell:~$
```

13.3 **tail**

Quite frequently, when we analyze log files, we are most interested into the latest events that have been logged. These events are normally written at the bottom of the log file. For example, in a server, there can be a log file that keeps tracks of all network requests. If a malicious request caused the system to shut down, you can verify which request caused the shut down by looking at the last entry in the log file.

For scenarios such as browsing log files, as discussed above, a tool called **tail** is made available. Its behaviour is nearly identical to that of **head**. The only difference is that, instead of showing the first ten (10) lines of the file, it shows the last ten (10) lines of it. Also, just like with the **head** command, you can specify the number of lines to display using the **-n** option followed by the number of lines you wish to display. See the example below:

```
marcel@dell:~$ tail -n 3 poem1
So Eden sank to grief,
So dawn goes down to day.
Nothing gold can stay.
marcel@dell:~$
```

If you want to look at log files that are currently being updated, you can use **tail** with the **-F** option. For example, **tail -F /var/log/syslog** will display the last ten lines of **syslog***, and update these lines if a new content is added to file. To stop monitoring the file, and get a command prompt back, you need to press **Ctrl+C**.

**This is a log file that collects messages from various programs and services, including the kernel.*

13.4 **tac**

In some scenarios, such as when analyzing some log files, you may want to display an entire file with the last lines appearing first. I.e., you may wish to read a particular file with its lines displayed in reverse order. For these scenarios, the **tac*** is made available. Like **cat**, **tac** can also be used to concatenate files. However, in the result, each individual file will appear with its lines in reverse order. See the example below:

Stands for **cat spelled backwards.*

```
marcel@dell:~$ tac poem2
By any other name would smell as sweet.
What's in a name? That which we call a rose
marcel@dell:~$
```

13.5 **sort**

Suppose you have data stored in a text file like the one shown below, where each line represents one element, and each column represents one category.

USA	321	9,826
China	1,367	9,596
UK	64	243
Brazil	204	8,511
Canada	35	9,984
India	1,251	3,287
Egypt	88	1,001

Listing 13.3: Countries file.

In this example, the first column represents the country's name, the second line its population (in millions), and the third line represents its area (in thousands of square kilometers). Each line, or row, represents one country.

It is clear that this file is not sorted in any obvious way. Not alphabetically, not by population, and neither by area. However, there is a tool called **sort** that can be used to sort files like this in any way you choose. The choices of how to sort the file, i.e., by which column, in which order, etc., are made by calling **sort** with the proper set of options. In Table 13.1, we present some of the most common options:

-k N	Sorts the file based on its <i>N</i> <i>th</i> column.
-r	Sorts in reverse order, i.e., from Z to A, in case of an alphabetical sorting, or from larger to smaller in case of a numerical sort.
-n	Sorts numerically.
-t	Chooses a column separator other than non-blank to blank. You need to enter the column separator inside double quotation marks

Table 13.1: **sort** options.

In the examples below, we show a number of sorting operations on the Listing 13.3. Note that, just like with **cat**, **sort** does not alter the

files provided as an argument. It simply displays the sorted version of it on your terminal. We will cover a method to save these alterations in Chapter XXX.

```
marcel@dell:~$ sort countries
Brazil      204      8,511
Canada      35       9,984
China       1,367    9,596
Egypt       88       1,001
India       1,251    3,287
UK          64       243
USA         321     9,826
marcel@dell:~$ sort -k2 -n countries
Canada 35 9,984
UK 64 243
Egypt 88 1,001
Brazil 204 8,511
USA 321 9,826
India 1,251 3,287
China 1,367 9,596
marcel@dell:~$ sort -k3 -n -r countries
Canada 35 9,984
USA 321 9,826
China 1,367 9,596
Brazil 204 8,511
India 1,251 3,287
Egypt 88 1,001
UK 64 243
```

In Listing 13.3, different columns (fields) are separated by tabs. By default, **sort** uses blank to non-blank transitions, i.e., transitions from spaces and tabs to visible characters, to mark the beginning of each column. In case we have a different separator*, we need to indicate the separator we are using with the option **-t** followed by the chosen separator written inside double quotes, as seen in the example below:

**Such as in
comma-separated
values, CSV, files.*

```
marcel@dell:~$ sort -t"," -k2 -n countries
Canada,35,9984
UK,64,243
Egypt,88,1001
Brazil,204,8511
USA,321,9826
India,1251,3287
China,1367,9596
```

Numerical versus Alphabetical sorting

When we want to sort a number of items alphabetically, we need only to compare the first character of each item. For example, we can use the fact that **Adam** starts with an **A**, and **Bryan** with a **B**, to sort these two items, regardless of the rest of them. In the cases where the first character of two different items is the same, we need to check the second character. For example, we can check that **Adam** comes before **Aidan**, because **d** comes before **i**. The same goes if the first two characters are the same, in which case the third item is compared, and so it goes.

When performing a numerical sort, we always need to look at the whole number. If we only look at the first character, we would have reached the conclusion that the number **9** is greater than **10**, because **9** is greater than **1**. This is clearly not what we want when sorting numbers. Hence, this is why **sort** needs to know if the user is trying to sort things numerically or alphabetically. By default, **sort** assumes that an alphabetical sort is taking place.

13.6 cut

When dealing with data stored in text files, we are often only interested in a subset of it. For example, in Listing 13.3, we could be interested only in the countries population, or only in their area, as opposed of both. For scenarios like these, **cut** comes in handy. As the name suggests, it is used to cut, or in other words remove, any contents from a text file that the user doesn't need.

In the example shown below, we first use this filter to remove the information about the countries' area. Following, we use it to remove all information about their population.

```
marcel@dell:~$ cut -f 1,2 countries
```

```
USA      321
China    1,367
UK        64
Brazil    204
Canada    35
India     1,251
Egypt     88
```

```
marcel@dell:~$ cut -f 1,3 countries
```

```
USA      9,826
China    9,596
UK        243
Brazil    8,511
Canada    9,984
```

```
India    3,287
Egypt    1,001
```

Note that, with **cut**, we need to specify the fields, or categories, that we wish to keep, not the ones we want to cut. This is done by using the **-f** option, followed by the column numbers separated by commas. If you are interested in a range of columns, for example columns one (1) to three (3), you can use a dash, as in **cut -f 1-3 countries**.

Note that, **cut** uses a tab as its default delimiter, not the transition from blank to non-blank characters like **sort**. If multiple tabs, or spaces and tabs occur between two columns, **cut** will probably not behave the way you expect. In case any other separator is used, for example, commas or colons, the user needs to specify the separator using the **-d** option followed by the desired separator, written between single quotes. See the example below for a **cut** command being applied to a CSV version of Listing 13.3.

```
marcel@dell:~$ cut -d',' -f 1,2 countries
USA,321
China,1,367
UK,64
Brazil,204
Canada,35
India,1,251
Egypt,88
```

13.7 sed

sed* can be used to find lines containing a particular word or [regex](#), similarly to **grep**, or to display all lines that do not contain a particular word or [regex](#), which can also be done using **grep**. However, there is one particular task for which **sed** excels, and for which it is arguably the best tool available: replacing particular words or expressions in a text file with other words or expressions*.

**Short for stream editor.*

sed works by parsing the provided text file line by line, and then making changes whenever there is a match to the provided word or [regex](#). The syntax to use **sed** to search and replace words or [regex](#) inside a text file is the following: **sed "s/original_word/replacement_word/" file_name**. See the example below:

**Like find and replace tools that exist in most text editors.*

```
marcel@dell:~$ cat poem
A rose is a rose is a rose
Gertrud Stein
marcel@dell:~$ sed "s/rose/flower/" poem
A flower is a rose is a rose
Gertrud Stein
```

Note that, in the example above, only the first occurrence of the word **rose** was replaced by **flower** on the first line. This happens

because, by default, **sed** skips to the next line after finding a match. In order to replace all words or [regex](#) in these scenarios, you need to specify that you want to perform a global search and replace, using the following syntax: **sed "s/original_word/replacement_word/g" file_name**.

```
marcel@dell:~$ cat poem
A rose is a rose is a rose
Gertrud Stein
marcel@dell:~$ sed "s/rose/flower/g" poem
A flower is a flower is a flower
Gertrud Stein
```

Note that, just like all other filters discussed in this chapter, **sed** does not alter the text file provided as an input. It will simply display its result in the terminal. In Chapter XX, we will discuss methods to save the alterations made by **sed**.

EXERCISES

For all questions below, you should use the following text files (copy and paste them into three text files called **poem1**, **poem2**, and **countries** respectively):

```
What's in a name? That which we call a rose
By any other name would smell as sweet.
```

```
A rose is a rose is a rose
```

USA	321	9,826
China	1,367	9,596
UK	64	243
Brazil	204	8,511
Canada	35	9,984
India	1,251	3,287
Egypt	88	1,001

1. Concatenate both poems, displaying the result on the terminal.
2. Concatenate both poems, but displaying each one of them in reverse order. I.e., each poem should start at its last line, and proceed all the way to its first.
3. Display only the first three lines of the file **countries**.
4. Display only the last three lines of the file **countries**.
5. Use **tail** with the proper options to keep track of the last five (5) lines of a log file called **sysgen.log** that is currently being updated.

6. Sort the **countries** file based on its population (second column) in crescent order.
7. Sort the **countries** file based on its area (third column) in descent order.
8. Use **sed** to display the contents of **poem2** with all instances of the word **rose** replaced by **flower**.
9. Does the **sed** command changes the file provided as an argument? For example, after running a **sed** command on a file called **poem2**, do the modifications performed by **sed** get stored in memory?

AWK, RENAME, AND FIND

In our previous chapter, we have covered a number of simple tools that could be applied to text files, such as **cat**, **sort**, **sed**, etc. In this chapter, we will cover some slightly more complex, but incredible powerful, tools that allow us to filter lines of a text file based on a chosen criteria (**awk**), rename multiple files at once (**rename**), and find specific files according to different search criteria (**find**).

14.1 **awk**

In previous chapters, we showed how to use **grep** to display lines of text that contained a particular word or [regex](#). However, what if we want to display all lines of a particular text file that satisfy a particular pattern (condition), instead of a [regex](#)? For example, see the text file below:

```
toyota corolla 1970 2500
chevy malibu 1999 3000
ford mustang 1965 10000
volvo s80 1998 9850
ford thundbd 2003 10500
chevy malibu 2000 3500
honda civic 1985 450
honda accord 2001 6000
ford taurus 2004 17000
toyota rav4 2002 750
chevy impala 1985 1550
ford explor 2003 9500
```

Listing 14.1: Car dealership text file.

Listing [14.1](#) contains a number of columns with information about cars at a car dealership. The first column contains the car model, the second column its manufacturer, the third its year, and the fourth its price.

Clearly, we can use **grep** to display all the lines in which the car manufacturer is **Toyota** or **Ford**. However, we cannot simply write a [regex](#) and use **grep** to display all lines in which the car is from 1999 or newer. Also, we wouldn't be able to use a [regex](#) to display all lines in which the cars' prices are less than 15,000 CADs. For these scenarios, another tool, called **awk**^{*} can be very useful.

^{*}The name of this command stands for the first letter of the name of its authors.

awk is a pattern-scanning and processing language designed for Unix, and later rewritten by the [GNU](#) project as **gawk***. It searches for patterns and/or conditions in text files and displays the lines in which they were found. For instance, using Listing 14.1, we could have retrieved all lines containing cars made during or after 1999 with the following syntax:

*See the **awk** vs **gawk** text box.

```
marcel@dell:~$ awk '$3 >= 1999' cars
chevy malibu 1999 3000
ford thundbd 2003 10500
chevy malibu 2000 3500
honda accord 2001 6000
ford taurus 2004 17000
toyota rav4 2002 750
ford explor 2003 9500
```

Also, we could have retrieved all lines containing cars with a price tag at or below 9,000 CADs, with the following syntax:

```
marcel@dell:~$ awk '$4 <= 9000' cars
toyota corolla 1970 2500
chevy malibu 1999 3000
chevy malibu 2000 3500
honda civic 1985 450
honda accord 2001 6000
toyota rav4 2002 750
chevy impala 1985 1550
```

Note that we use the notation **\$N**, where **N** denotes the number of the column we want to check for an specific pattern. Also, note that we should always put the complete pattern statement within single quotes. Some of the most important patterns that can be used with **awk** are described in Table 14.1.

<	Less than
<=	Less or equal than
==	Equal to
!=	Not equal to
>	Greater than

Table 14.1: **awk** patterns.

Note that nothing prevents us from applying multiple patterns at once. For example, if we wanted to display all lines containing cars made during or after 1999 with a price tag at or below 15,000 CADs, we could have used the following syntax:

```
marcel@dell:~$ awk '$4 <= 15000 && $3 >= 1999' cars
chevy malibu 1999 3000
```

```
chevy malibu 2000 3500
honda accord 2001 6000
toyota rav4 2002 750
```

Where the **&&** operator denotes a logical **and**. I.e., it only returns the lines in which both conditions are met. If we wanted to display all lines in which either one condition **or** another condition were met (or both), we could have used the **||** operator.

Before we conclude this section, it is worth to note that **awk** is capable of doing much more than what we have described here. However, most of its other uses can also be carried out by other commands we have previously discussed such as **grep**, **cut**, or by combining commands, as we will discuss in Chapter XXX.

awk vs gawk

In this book, we have been using tools such as **ls**, and **cat**, that were rewritten by the GNU Project to replace Unix tools with the same name. For example, when we discussed **cat**, we discussed the GNU Project implementation of the **cat** tool originally written for Unix. However, when the **awk** tool was rewritten, the GNU Project developers added a **g** to its name, and then called it **gawk**.

As a result of this name change, some Linux distributions nowadays require users to write **gawk** to use this tool, whereas other distributions require users to write **awk**. Some Linux distributions even allow users to write **gawk** or **awk** to call the same tool.

14.2 rename

We have seen in Chapter 6 that we can use **mv** to rename individual files. For example, **mv original_name new_name** changes a file name from **original_name** to **new_name**. However, what if we want to change the name of multiple files at once? For example, what if we want to change the ending of all **.html** files from a website to **.php** so that they could be processed by a server-side script prior to be delivered to the end-user? For such scenarios, **rename** can be a very helpful tool.

rename uses [regex](#) to match patterns for filenames, and substitute it by a string literal. It also uses wildcards to indicate in which files the pattern matching should be applied upon. For instance, to change the ending of all **.log** files in a directory to **.bak**, we could have used the following command*:

```
marcel@dell:website$ ls
index.html contact.html about.html
marcel@dell:website$ rename 's /\.html$/\.php/' *
```

**The syntax presented below only works in Debian based systems, such as Ubuntu.*

```
marcel@dell:website$ ls
index.php contact.php about.php
```

Note that **rename** uses a syntax very similar to **sed**, discussed in the previous chapter. I.e., it is called with:

```
rename 's/original\_regex/replacement\_string/' wildcard
```

Also, note that `\.` is used to indicate that we want to use the period as a literal, as opposed to be a [regex](#) that can stand for any character.

As another example, we might have wanted to delete the ending of `.log` files, as opposed to switched it to something else. This could have been accomplished with:

```
marcel@dell:log$ ls
netlist.log tcpdump.log read_me
marcel@dell:log$ rename 's/\.log$//' *.log
marcel@dell:log$ ls
netlist tcpdump read_me
```

14.3 find

As the name suggests, **find** can be used to find files based on a number of possible search criteria. For example, you can search all files in a given directory whose name matches a particular wildcard pattern*. As another example, you can use **find** to search for all files in your system that were created after a certain date. You can even use this tool to search for all files that belong to a particular user, or have a particular set of permissions, among other possibilities. The syntax for the **find** is:

**Note that we need to use wildcards, not [regex](#) with this tool.*

```
find target\_directory criteria\_type criteria\_match
```

In case **criteria_match** uses wildcards, it needs to be placed within single quotes.

On table [14.2](#), you can find a list of different types of searches that can be done using **find**. Note that this list is not complete. For a complete list, refer to **find**'s man page.

Below we have a few examples of **find** being used in different forms:

```
marcel@dell:~$ ls -l
-rw-rw-r-- 1 marcel marcel    125  Sep 18 10:52 assign.pdf
-rw-rw-r-- 1 john  john     3212  Sep 23 21:51 classes.docx
-rw-rw-r-- 1 john  john      450  Sep 15 14:22 test1.pdf
-rw-rw-r-- 1 marcel marcel   1232  Jul 19 12:44 honesty.pdf
marcel@dell:~$ find . -name '*.pdf'
.
./assign.pdf
./test1.pdf
```

```

./honesty.pdf
marcel@dell:~$ find . -user john
.
./classes.docx
./test1.pdf
marcel@dell:~$ find . -mtime -10 #today is September 23
.
./assign.pdf
./classes.docx
./test1.pdf
marcel@dell:~$ find . -maxdepth 1 -size +1k
.
./classes.docx
./honesty.pdf

```

By default, **find** searches not only the provided directory, but also all its subfolders. You can change this behaviour by using the **-maxdepth** option with a value of **1**. You can also use it to perform searches on the current working directory and its immediate subfolders, by setting **-maxdepth** to **2**, and so on.

-name	Searches for files that match a provided wildcard.
-size	Searches for files whose size is smaller, equal, or greater than a provided number.
-type	Searches for all files of a particular type. Possible types are: regular files (-f), directories (-d), symbolic links (-l), among others.
-user	Searches for all files that belong to a particular user.
-samefile	Searches for all files with the same inode number as a file provided as an argument. In other words, it finds all hard links of a particular file.

Table 14.2: **find** search criteria types.

EXERCISES

1. Which command can be used to list all cars in Listing 14.1 that were built after 2000?
2. Which command can be used to list all cars in 14.1 that cost less than 15000 CADs?
3. Which command can be used to list all cars in 14.1 that were built after 2000 and cost less than 15000 CADs?
4. Which command can be used to change the ending of all **.jpeg** files in the current working directory to **.jpg**?
5. Which command can be used to change the ending of all **.jpeg** files in the current working directory to **.jpg**?
6. Which command can be used to change all uppercase letters, in all files in the current working directory, to lowercase?
*hint: check **rename**'s man page.*
7. Write down a command to search all files with more than 20kb inside the current working directory, as well as inside its immediate subfolders (but not on subfolders of subfolders).
8. Write down a command to search all files that have been modified in the past week inside the current working directory, as well as in all its subfolders.
9. Write down a command to search all files that have been accessed after a file named **syslog** has been modified.
10. Write down a command to search for all hard links of a file called **file**, that exists in your current working directory, in your entire file system.

REDIRECTION AND PIPING

In Chapters 13 and 14, the results of all examples were displayed on the terminal. I.e., no permanent changes were made to the text files provided as arguments, or any other files in the system. Also, we have worked with each tool by itself. I.e., we did not try to combine tools together, as according to the Unix philosophy.

In this chapter, we will address these two issues. First, we will show how to use redirection to save outputs in text files, as opposed to displaying them in the terminal. Following, we will cover how to combine two or more tools together using the concept of piping.

15.1 REDIRECTION

The concept of redirection is connected to that of data streams, which are the the source of system inputs, as well as the destination of system outputs.

There are three data streams in **bash**, one input, and two outputs, which are listed on Table 15.1.

stdin	The standard input, as the name suggests, refers to the source of the data that is being fed into the shell. It is normally linked to the keyboard.
stdout	The standard output, as the name suggests, refers to where the output of any successful commands should go. By default, it is linked to the terminal.
stderr	The standard error stream refers to where the output of unsuccessful commands should go. Like stdout , it is also linked to the terminal by default.

Table 15.1: Bash data streams.

In Figure 15.1, you can see a depiction of these three data streams, as well as their relationship with the shell.

Even though both output streams are linked to the terminal display by default, separating them has one great advantage. It allow users to deal with the output of commands in different ways, depending if they were succesful or not. For example, you can choose to save




Figure 15.1: Visual representation of data streams.

the output of successful commands to a particular file, but send error messages to **/dev/null***

**See /dev/null box below.*

/dev/null

In Linux systems, **/dev/null** is a simple device implemented in software, which does not correspond to any hardware device. Simply put, it is used to throw information away, or to test reading with empty information. In a nutshell:

- **/dev/null** looks as an empty file when you try to read from it using commands such as **cat** or **less**.
- Redirecting data to **/dev/null** results in that data being thrown away. It simply “disappears” from the system.

15.1.1 Redirection Syntax

stdout

To redirect the standard output, **stdout**, to a given file, you need just to add **> file_name**, or **1> file_name**, to the end of the command. In the example below, which uses the file **countries** presented in Listing 13.3, we save the output of a sorting operation into a file called **countries_sorted**.

```
marcel@dell:~$ sort countries > countries\_sorted
marcel@dell:~$ cat countries\_sorted
Brazil    204    8,511
Canada    35     9,984
China     1,367  9,596
Egypt     88     1,001
India     1,251  3,287
UK         64     243
USA       321    9,826
```

It is important to note that because we are redirecting the output of the sorting operation, we don’t get any feedback on our terminal. If we want to see the result, we need to use a command to read the **countries_sorted** file. In the example above, we used **cat**. Also, note that if the file **countries_sorted** already exists, it is overwritten. In

case the command results in an error, the file **countries_sorted** will be empty.

In fact, when applying redirection, the first action taken by the shell is to erase any data from the destination file. This can lead, sometimes, to an unexpected behaviour. For example, a command **sort -k2 countries > countries** will result in an empty file. This happens because the file **countries** will be emptied prior to the sorting operation.

stderr

To redirect the standard error, **stderr**, to a given file, you need just to add **2> file_name**, to the end of the command.

You can also redirect both output streams to the same file by adding **&> file_name**, or you can redirect each stream to an specified file using **1> file_name1 2> file_name2** at the end of the command.

stdin

The standard input normally comes from the keyboard, in the form of arguments that we provide. For example, **ls folder_name** uses the **stdin** to provide the argument **folder_name** to the **ls** command.

More often than not, commands that are regularly applied to files accept both the file name, or a redirection from the file. For example, both commands below are equivalent:

```
marcel@dell:~$ head -n3 countries\_sorted
Brazil    204      8,511
Canada    35        9,984
China     1,367     9,596
marcel@dell:~$ head -n3 < countries\_sorted
Brazil    204      8,511
Canada    35        9,984
China     1,367     9,596
```

Because of equivalences such as the one showed in the example above, scenarios in which you need to redirect the **stdin** are not as common as scenarios in which **stderr** and **stdout** are redirected.

15.2 PIPING

The concept of piping commands together is quite simple. It can be summarized as: “*using the output of one command as an input to another command*”. In fact, the name piping derives from an analogy with a real world pipeline, where the output of each pipe becomes the input of another pipe further down the pipeline.

Regardless of its simplicity, piping is a very powerful concept. To better understand its implications, we will discuss in what follows a scenario in which piping commands together solves a particular problem.

15.2.1 Piping Application Scenario

A college keeps the records of all their students in separate files, with each file corresponding to the programs they are enrolled such as the ones shown in Listings 15.1 and 15.2*. Now, suppose that you are requested to come up with a list of the top three students in the entire college, based on their GPAs.

**In a real world scenario, these files would be much bigger.*

```
John Smith john.smith@mycollege.ca 2.9
Jane Doe jane.doe@mycollege.ca 3.9
Mohammad Ali mohammad.ali@mycollege.ca 3.4
Chun Li chun.li@mycollege.ca 3.1
```

Listing 15.1: Computing Science students information (cs.info).

```
Aditya Kapoor aditya.kapoor@mycollege.ca 3.9
Juan Sanchez juan.sanchez@mycollege.ca 2.8
Klaus Klein klaus.klein@mycollege.ca 3.6
Marc Belleville marc.beleville@mycollege.ca 2.6
```

Listing 15.2: Fine Arts students information (fa.info).

Clearly, this scenario could use a tool that combines the power of **cat** to concatenate text files, **cut** to get rid of unnecessary columns, and **sort** to sort the students according to their GPAs (as well as **head** or **tail**, to limit the output to only three students).

With piping, the solution to this problem is quite simple. We first use **cat** to join all text files with students information. However, instead of displaying the results on the terminal, or redirect it to a file, we pipe it to **cut** which uses a space as its delimiter and keeps only the first, second, and fourth fields. Following, we send the output of **cut** to **sort**, which should sort students using its third column,* numerically, and in descending order. Finally, we pipe the output of **sort** to **head**, which is used to display only the first three lines of the result, as seen below:

Note that the output of the **cut, which is the input of **sort**, has only three columns.*

```
marcel@dell:~$ cat cs.info fa.info | cut -d" " -f 1,2,4 |
    sort -r -n -k3 | head -n3
Jane Doe 3.9
Aditya Kapoor 3.9
```

Klaus Klein 3.6

marcel@dell:~\$

15.2.2 Piping Syntax

To pipe the output of one command as an input to another command, as shown in the example above, we use the piping operator `|`.

```
marcel@dell:~$ command1 options input_arguments | command2
options | command3 options | ...
```

Note that there is no limit on the number of pipes that can be used. Also, note that we exclude the input arguments* for all pipelined commands besides the first. This happens because, for these commands, the input should come from the pipe, not from another file. In other words, the input coming from the pipe replaces the input arguments we normally use.

**Which are normally file names.*

15.2.3 Piping examples

The application scenario we presented had a quite complex pipeline of commands. To get a better grasp of the concept of piping, see the examples below, where simpler pipelined commands are used:

```
marcel@dell:ls | less
```

In this example, we simply redirect the list of files and folders in the current directory to be displayed in **less**, as opposed to be printed in the terminal. This is a very common use of pipes, as it allows users to check the contents of large directories that would not fit the terminal screen.

```
marcel@dell:cat *.log | grep "error" | less
```

In this example, we first concatenate all files with a **.log** ending, and then we use **grep** to find the lines in which the string **error** can be found. Finally, the result is displayed in **less**.

```
marcel@dell:find / -user john | xargs grep "password"
```

In this example, we first use **find** to get the name of all files that belong to a user called **john**. Then, we use **grep** to search for the string **password** inside all of these files. Note that we used a command called **xargs*** between the pipe and **grep**. This was done because we wanted **grep** to search inside the files whose names and paths were passed by the pipe. Without **xargs**, **grep** would perform its search on the list of filenames, as opposed to performing the search inside all files presented by this list.

The **xargs command breaks the list of arguments passed by the pipe into its individual elements.*

15.2.4 *Piping and Redirection*

Just like we did for regular commands, you can redirect the final result of a pipeline into a text file using the **>** operator.

In all examples from the previous section, we could have written the results into a text file called **file**, as opposed to displaying it in **less** or printing it in the terminal, by using:

```
marcel@dell:ls > file
marcel@dell:cat *.log | grep "error" > file
marcel@dell:find / -user john | xargs grep "password" > file
```

15.2.5 **tee**

In our previous section, we showed how to redirect the final result of a pipeline. However, we cannot simply redirect the output of a pipeline midway. For example, the command below will result in an error.

```
marcel@dell:cat *.log > my\_file | grep "error" | less
```

This should come to no surprise, if we think about our pipeline real world analogy. If we dump the contents of a pipe, we cannot expect it to proceed further down the pipeline.

In order to be able to save intermediate results, we need to use **tee***. This tool allow us to copy the output of a pipe into a file, but also send it further down the pipeline. For example, we can fix the previous example by using **tee** to send the output of **cat** to a file called **my_file**, but also send it further down the pipeline, using the following syntax:

```
marcel@dell:cat *.log | tee my\_file | grep "error" | less
```

The name of this tool derives from the format of the letter **t.*

Part IV

SCRIPTING

In the following chapters, we introduce the concept of scripting, i.e., creating files containing a sequence of instructions that perform a specific task. Scripts can be used for automating repetitive tasks. They are widely used by system administrators to take care of tasks such as creating hundreds of new users, or configuring multiple machines. With scripts, such tasks can be done in a matter of minutes, instead of taking hours or days of repetitive work.

INTRODUCTION TO SCRIPTING

In movies and plays, a script is a document containing the lines that the actors and actresses should say, as well as the actions they should perform, i.e., running, jumping, fighting, etc. Similarly, in computing, a script is a text file containing a sequence of lines of code that we issue to a system in order for it to print outputs and execute particular actions.

In this chapter, we will start by introducing a very basic script. Then, we explain how to run scripts using the terminal. Following, we discuss the concepts of variables and shell environments. Finally, we end this chapter by showing how a script can work with inputs provided by the user in real-time.

16.1 HELLO WORLD SCRIPT

Most programming languages books start by introducing a source code to print the sentence "Hello World" on the terminal. This book is no exception, in Listing 16.1 we show a very basic **bash** script* that, once executed, prints this sentence.

```
1 #!/bin/bash
2 #Hello World example
3
4 echo "Hello World!"
```

Listing 16.1: Script for a Hello World example.

There is a discussion about the meaning of the words **script and **source code** in the Compile languages vs Interpreted languages box on page 105.*

Even though this is a very simple script, it already contains a number of basic features that will be present in almost all **bash** scripts we are going to be writing in the course of this book. These features are discussed in what follows.

16.1.1 **shebang** line

The first line of any terminal script, no matter which language they were written with, must tell the shell which interpreter to use in order to read the script. This line is usually called a **shebang** line.

A **shebang** line starts with a pound sign (hashtag), **#**, followed by an exclamation sign, **!**, followed by the location of the tool used to interpret the script. Because we are writing **bash** scripts in this book, we need to provide the location of **bash**. In most distributions, **bash**

is at the **/bin** folder*. Hence, in the first line of our script we have a shebang line: **#!/bin/bash**.

You can confirm **bash's location in different distributions by issuing a **which bash** command.*

16.1.2 Comments

Many times we want to write some information into a script that we don't want the interpreter to process. For example, you may want to write the name of the author of the script, or provide a brief description of what it does. In both cases, the information you are providing is just for yourself, or other users, to read. The interpreter should simply ignore these lines.

For scenarios such as the ones described above, we can make use of comments. A comment is a block of text that follows a pound sign (hashtag) **#**. With the sole exception of the **shebang** line we discussed in the previous section, the interpreter ignores everything written between a **#** sign and the end of the line. In Listing 16.1, we have a comment on the second line. Note that in our example, we wrote a comment on a blank line. However, nothing prevents us from adding a comment to the right of a non-empty line.

In large scripts, comments are often used to describe specific parts of it. They can be very helpful for others, and even your future self, to understand the script. It is a good practice, in such cases, to explain what is it that the script does, not how it does it.

16.1.3 **echo**

echo prints on the terminal the arguments passed to it. In Listing 16.1, we use it in line 4 to print the string "Hello World!".

Note that whenever we want to print a string, i.e., a sequence of words, we need to write it inside quotes. There are two types of quotes that can be used:

SINGLE QUOTES Also known as hard quotes, allows **echo** to print what is inside the quotes in verbatim. I.e., exactly as it is written.

DOUBLE QUOTES Allow **echo** to replace the name of variables preceded by a dollar sign **\$** with the value stored in it. We will cover variables later in this chapter.

Compiled languages vs Interpreted languages

There is a plethora of computer languages available today, each one catering for a different kind of public, or different kind of application scenarios. There are multiple ways to categorize these languages. However, one of the most fundamental di-

vides between computer languages is that some are compiled languages, and others are interpreted languages.

A compiled language requires a tool called compiler to translate a **source code**, which is written by a software developer, into an executable file. Then, this executable file can be easily called by the user to perform whatever tasks it was designed to do. If the developer needs to change the application, he or she will need to alter the original source code and recompile it. Note that, once the executable file has been created, the system does not need the source code anymore. **C**, **C++**, and **Java**, are some examples of compiled languages.

An interpreted language, on the other hand, requires a tool called interpreter. This tool reads the script line by line, processing all commands it finds as it goes through the script. **Bash**, **Python**, **PHP**, and **JavaScript** are some examples of interpreted languages.

Note that in the paragraphs above, I am calling the text files that serve as input to compilers **source codes**, and files that serve as input to interpreters **scripts**. This distinction is applied by this book's author, but is not universal. Some developers might talk about **source codes** for Python applications. That said, you will rarely see the expression **bash source code** being used instead of **bash script**.

16.2 USING THE TERMINAL TO RUN SCRIPTS

To run a **script**, you simply have to call it by its name using the terminal so that the shell can execute it. In fact, most of the tools this book has covered so far, such as **ls**, **cat**, **awk**, etc., are nothing but binary scripts that are located at the **/bin** directory.

Clearly, in order to execute scripts, the shell needs first to know where to find them. That is where a variable called **PATH** plays an important role, as we discuss in what follows:

16.2.1 *PATH* variable

A variable called **PATH** is used to store the paths of all directories that **bash** searches while it looks for scripts to execute. By default, in an Ubuntu distribution, it contains:

```
marcel@dell:~$ echo $PATH
/home/marcel/bin:/home/marcel/.local/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/usr/games:/usr/local/games:/snap/bin
```

Note that in your own system, your username should replace mine (**marcel**) in the `PATH` folders that are under the user's home folder.

When a user enters any command, the shell searches through all these directories, one by one, while trying to find a script which name matches the provided command. Once a match is found, the search is interrupted and the script is executed*.

It is important to note that the current folder is not usually in the `path`. However, we can always execute a script by providing the shell with its path. This means that, in case we want the shell to execute a script with name **script_name.sh** within the current folder*, we use the following syntax:

```
marcel@dell:~$ ./script_name.sh
```

Note that we could have used a similar notation to run scripts in any other directories by providing the shell with its relative or absolute paths.

This means that if two scripts with the same name exist in different directories in the **PATH, the one that comes first in the **PATH** is the one that gets executed.*

**Which is always indicated by a period sign (.).*

16.2.2 Granting permission to execute a Script

Suppose that you save the script provided in Listing 16.1 in a text file called **script.sh*** in your working directory. You should expect to be able to run it by simply writing down the command: **./script.sh**. However, when you do so, you get the error message shown below:

```
marcel@dell:~$ ./script_name.sh
bash: ./script.sh: Permission denied
```

What happens is that, by default, files created in most linux systems are created with only reading and writing permissions, not with executing permissions*. To change this, you need to grant yourself permission to run your script. This can be accomplished by issuing the command: **chmod +x script_name**, as shown below:

```
marcel@dell:~$ ./script_name.sh
bash: ./script.sh: Permission denied
marcel@dell:~$ chmod +x script.sh
marcel@dell:~$ ./script.sh
hello world
```

Note that, once we granted the script permission to be executed, we were able to call it without any problems. As a result, it printed **Hello World** on the terminal as expected, and exited.

It is important to note that, apart from the fact that it requires executing permission, a script is no different than any other text file. In fact, just like any other text file, it can be opened for reading using **more**, **less**, **cat**, or opened for editing in **nano**, **vim** or even with **GUI** tools such as **Gedit**.

Ending **bash script names with **.sh** is not necessary, but it is a good convention to follow.*

**A comprehensive discussion on file permissions can be seen in Chapter XXX.*

16.3 VARIABLES AND ENVIRONMENTS

A variable is a name we assign to a place in memory that we can use to store a particular value. For example, if we enter **var=6**, in our terminal, the shell saves the value **6** in memory on a variable called **var**. To retrieve the contents of a variable, we need only to call its name with a preceding dollar sign, **\$**. See the example below:

```
marcel@dell:~$var=6
marcel@dell:~$echo "The value stored in var is $var"
The value stored in var is 6
marcel@dell:~$var=8
marcel@dell:~$echo "The value stored in var is $var"
The value stored in var is 8
```

As shown above, a variable can have its value changed after it was first defined. Hence the name variable. Note that we used double quotes in the example above, instead of single quotes. This was done to ensure that the shell would translate **\$var** into the value stored in **var**, before printing the message. In case we had used single quotes, the result would have been quite different, as shown below:

```
marcel@dell:~$var=6
marcel@dell:~$echo 'The value stored in var is $var'
The value stored in var is $var
```

Variables are extensively used when working with scripts. They can be used to store numbers, words or sentences*, and even filenames.

When a user logs into the system, the shell normally sets up an environment for this specific user with number of variables that control the behaviour of the terminal. The **PATH** variable we discussed before is an example of a variable that is normally set with default values. Other examples are **LOGUSER**, and **HOSTNAME**, which store the user name of the user that is currently logged, and the host name, respectively. You can see all default variables by issuing a **env** command.

It is important to note that each user that logs into the system is assigned his or her own exclusive environment. Also, it is important to know that a variable is only defined within its own environment. This guarantees that multiple users can access the system at the same time without the risk of one user altering another user's variables.

**Which in computing jargon are called strings.*

16.3.1 **source**

When a user calls a script, the shell creates a separate environment to store all variables created by the script. I.e., the shell executes the script within its own environment*. Executing scripts in their own environments results in two important features:

**You can think about environments as sandboxes.*

- A script does not have access to any variables set prior to the start of the script by the user, or by the system.
- Any variables set by the script will cease to exist as soon as the script finishes execution.

By creating an exclusive environment for each script that is executed, the shell guarantees that scripts will not accidentally overwrite any variables that exist outside them. For example, if a script defines a variable called **PATH**, this variable will have no relationship to the default **PATH** variable defined by the shell. Furthermore, once the script has finished execution, the value stored in the **PATH** will be the same value it had before the script was executed. The example below demonstrates this behaviour for the **USER** variable. It calls a script called **change_user.sh**, shown in Listing 16.2.

```
1 #!/bin/bash
2 #script that changes the LOGNAME variable
3 LOGNAME=Jar
4 echo "The current username is: $LOGNAME"
```

Listing 16.2: Script for changing the **USER** variable.

```
marcel@dell:~$echo "The current username is: $LOGNAME"
The current username is: marcel
marcel@dell:~$./script.sh
The current username is: Jar
marcel@dell:~$echo "The current username is: $LOGNAME"
The current username is: marcel
```

It is possible to run scripts without creating an exclusive environment for them. I.e., running scripts in the same environment as the shell itself. To do so, you simply need to use the **source** command in front of the script call, using the syntax: **source ./script_name.sh**. This is called sourcing the script. An obvious scenario for which sourcing a script is useful is the one in which a user wants to alter or use shell variables from within the script. See the example below:

```
marcel@dell:~$echo "The current username is: $LOGNAME"
The current username is: marcel
marcel@dell:~$source ./script.sh
The current username is: Jar
marcel@dell:~$echo "The current username is: $LOGNAME"
The current username is: Jar
```

**Assuming the script is in your current working directory.*

16.4 GATHERING USER'S INPUT

There are two ways in which we can gather inputs from users that are calling a script. Using the **read** keyword, or passing user input as arguments. Both methods are discussed in what follows.

16.4.1 *read*

The **read** keyword, as the name suggests, reads an input from the user and saves it into a variable. Its syntax is: **read variable_name**. Note that the user can enter numbers, words, or even full sentences. See the example below:

```
marcel@dell:~$ ./read_input.sh
Enter your name:
Marcel
Hello Marcel
marcel@dell:~$ ./read_input.sh
Enter your name:
Marcel Jar
Hello Marcel Jar
```

This example uses the script **read_input.sh** in Listing 16.3, shown below:

```
1 #!/bin/bash
2 echo "Enter your name:"
3 read name
4 echo "Hello $name"
```

Listing 16.3: Script that uses the **read** keyword.

16.4.2 *Passing arguments to a script*

You can also gather inputs from users by having them passing arguments to the script during its call. This is done by writing all arguments, separated by spaces, when the script is called. The value of each argument can be retrieved from inside the script using the syntax **\$N**, where **N** stands for the argument number. I.e., if it is the first argument, the second, etc. You can also get the total number of arguments passed to the script using the syntax **\$#**. In the example below, we pass three arguments to a script call:

```
marcel@dell:~$ ./read_arguments.sh Marcel Jar 25
Hello Marcel Jar you are 25 years old
I was given 3 arguments.
```

This example uses the script **read_arguments.sh** in Listing 16.4.

```

1 #!/bin/bash
2 echo "Hello $1 $2 you are $3 years old"
3 echo "I was given $# arguments."

```

Listing 16.4: Script that takes arguments from the user's script call.

EXERCISES

1. Explain, with your own words, what is the purpose of a **shebang** line. Also, write the proper shebang lines for **Bash**, **Python** and **Pearl** scripts. *hint: use the **which** tool to find the location of the proper script interpreter.*
2. Write a **bash** script that simply prints on your screen a message: **This is my first script!**. Make sure to include a **shebang** line, and one comment line containing your name as the script's author. Also, remember that you need to grant the script permission to be executed with: **chmod +x script_name**.
3. Using your own words, explain what is a variable?
4. How can you create a variable called **name** and assign to it your own name using **bash**.
5. How can you create a variable called **age** and assign to it your own age?
6. Explain, with your own words, the advantage of having the shell creating a different environment for each user that logs into the system. Also, explain the advantage of creating, by default, separate environments for scripts, when the user calls them.
7. Explain, with your own words, what is the difference between calling a script with or without the **source** command.
8. Write a script that asks the user to enter its name, its job, and its age. Following, your script should print: **Hi USERNAME, you are a USERJOB, and you are USERAGE years old.**, where the values of **USERNAME**, **USERJOB**, and **USERAGE** are taken from the user's inputs. *Hint: Use the **read** keyword*
9. Rewrite the script from the previous question. In this new version, instead of using **read** to gather the user's input, the script takes three arguments during its call: **USERNAME**, **USERJOB**, and **USERAGE**, respectively.